Maximilian Schwarzmüller

Developer, AWS Solutions Architect, Online Instructor

academind.com

maximilian-schwarzmueller.com

@maxedapps

What exactly is “Generative AI”?

Generative AI?

Generative AI refers to a category of **artificial intelligence** models that are capable of **creating new content**—such as text, images, audio, video, or code—that is **similar to what a human might produce.**

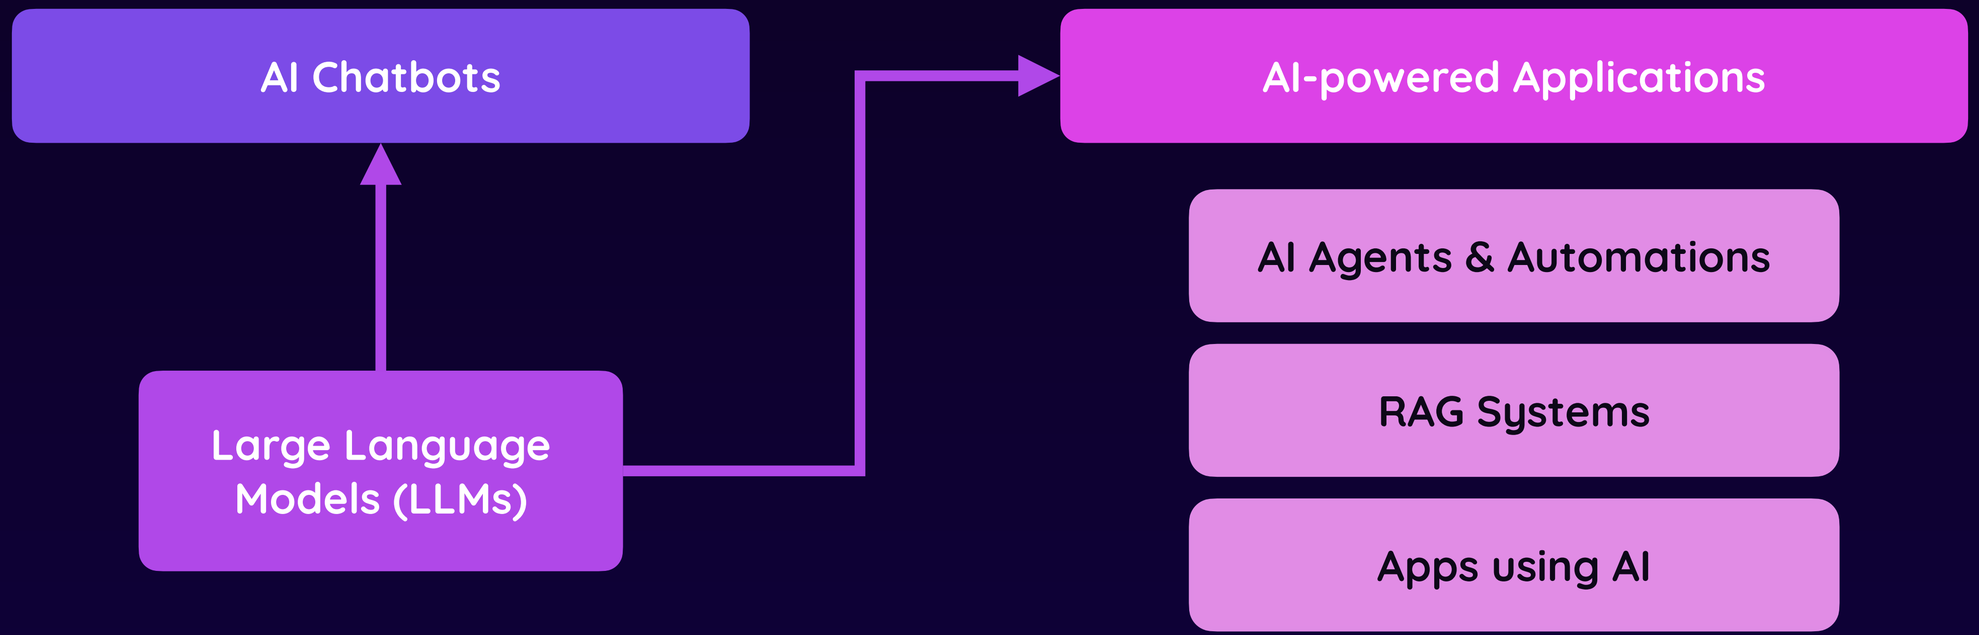

About This Course

AI-powered Apps & Workflows

- RAG, CAG & Finetuning
- Agents & Automations
- Custom AI Applications

Efficient Usage

- ChatGPT & Co Features
- Prompt Engineering
- Examples & Use-cases

Foundation

- Using AI Chatbots & Models
- Understanding Generative AI & LLMs

The Generative AI Revolution

Models	**ChatGPT** GPT-3	**ChatGPT** GPT-4	**Google Gemini**	**Anthropic Claude**	**ChatGPT** GPT-4o	**DeepSeek R1**	. . .
Capabilities	**Images**	**Code Execution**	**Audio**	**Web Search**	**Deep Research**	**Agents**	. . .

Highly Dynamic: New AI models & services built on top of those models are launched all the time

Highly Innovative: New application areas and “tools” are unlocked frequently

There Are Multiple Ways Of Using Gen-AI

Services / Apps

Service by third-party provider, used through their app / site

- **AI Chatbots**
  ChatGPT, Gemini, Grok, ...
- **Product Features**
  Photoshop, Cursor, Excel, ...

Main focus of this course

Integrated via APIs

API by third-party provider, integrated into your app

- **OpenAI & Co APIs**
  Exposed via REST & SDKs

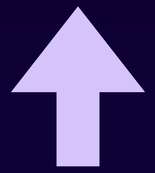

Also covered in-depth, including how to build custom apps & agents

Self-Hosted

Models hosted on your (rented) hardware

- **Open Models**
  LLama, DeepSeek, ...
- **Self-trained Models**
  If you have the resources...

Self-hosting intro included, training your own models is NOT covered

Key Gen-AI Service Providers—Overview

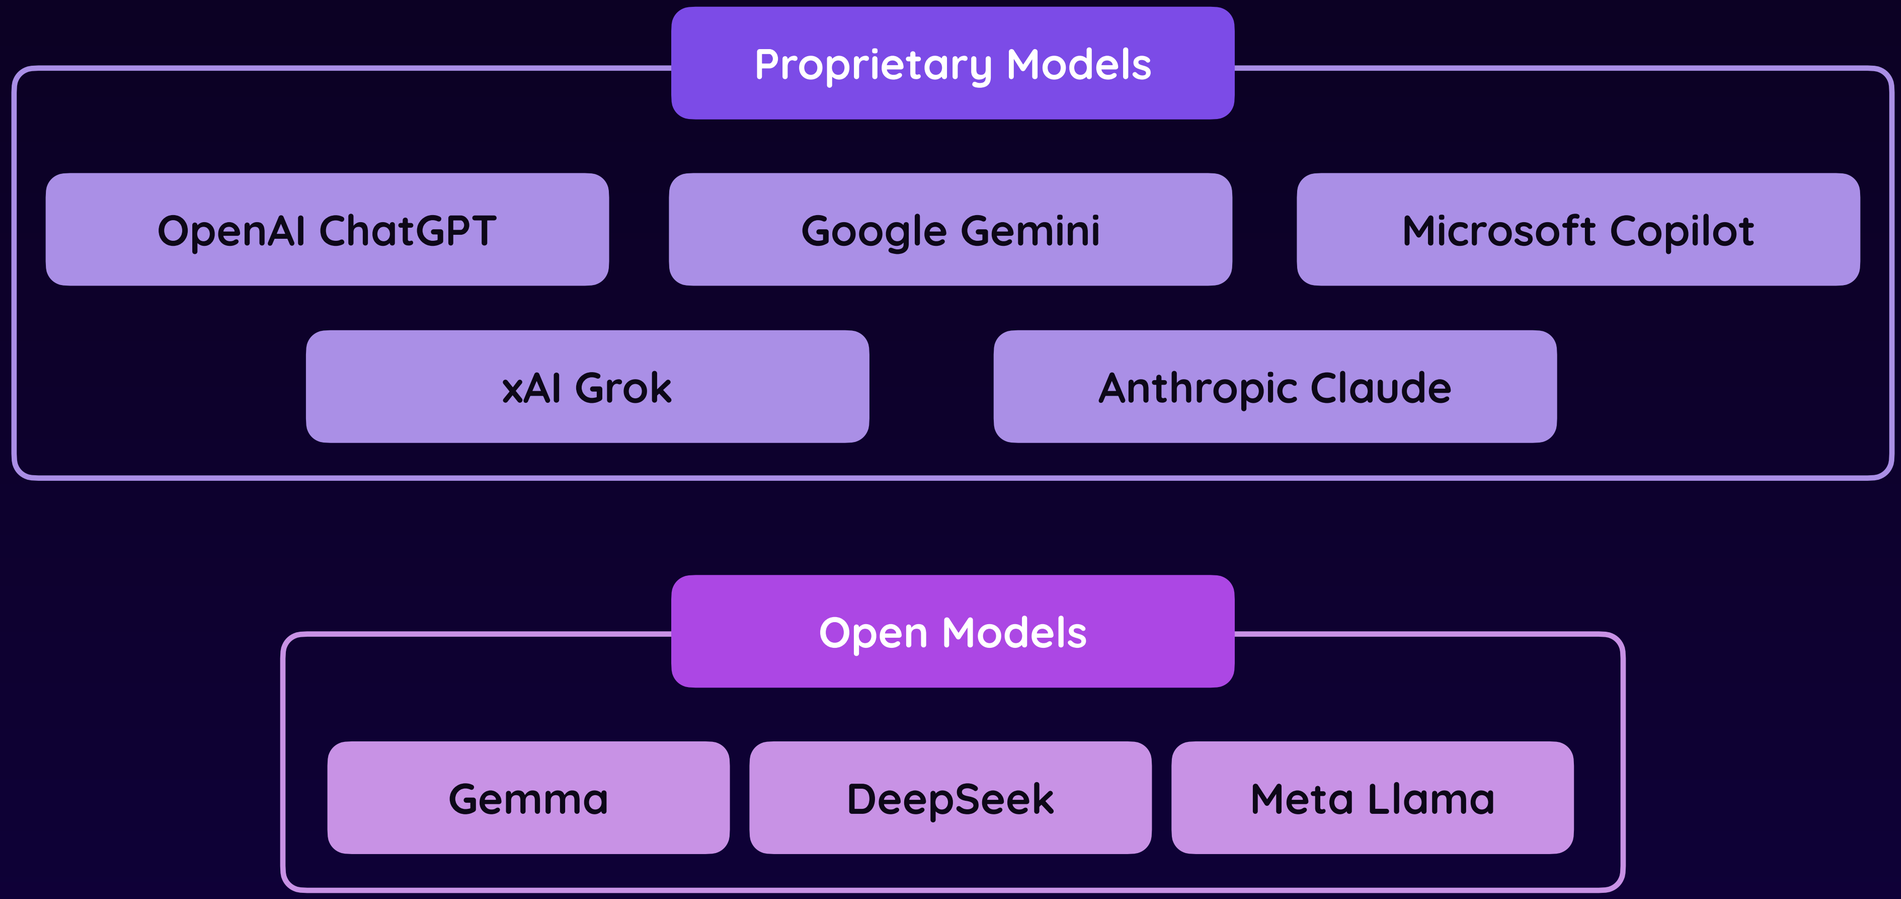

Important

You will not necessarily be able to reproduce the responses you see in the videos!

Even when using the exact same prompts.
All these AI models have a certain degree of “randomness”
& the underlying models will evolve and change.

Important

Free Usage Is Possible

But often not enough

ACADEMIND

Generative AI — Behind The Scenes

How LLMs Works Under The Hood

- How ChatGPT & LLMs Were Trained
- How They Work
- Features & Limitations

ChatGPT & Co Under The Hood

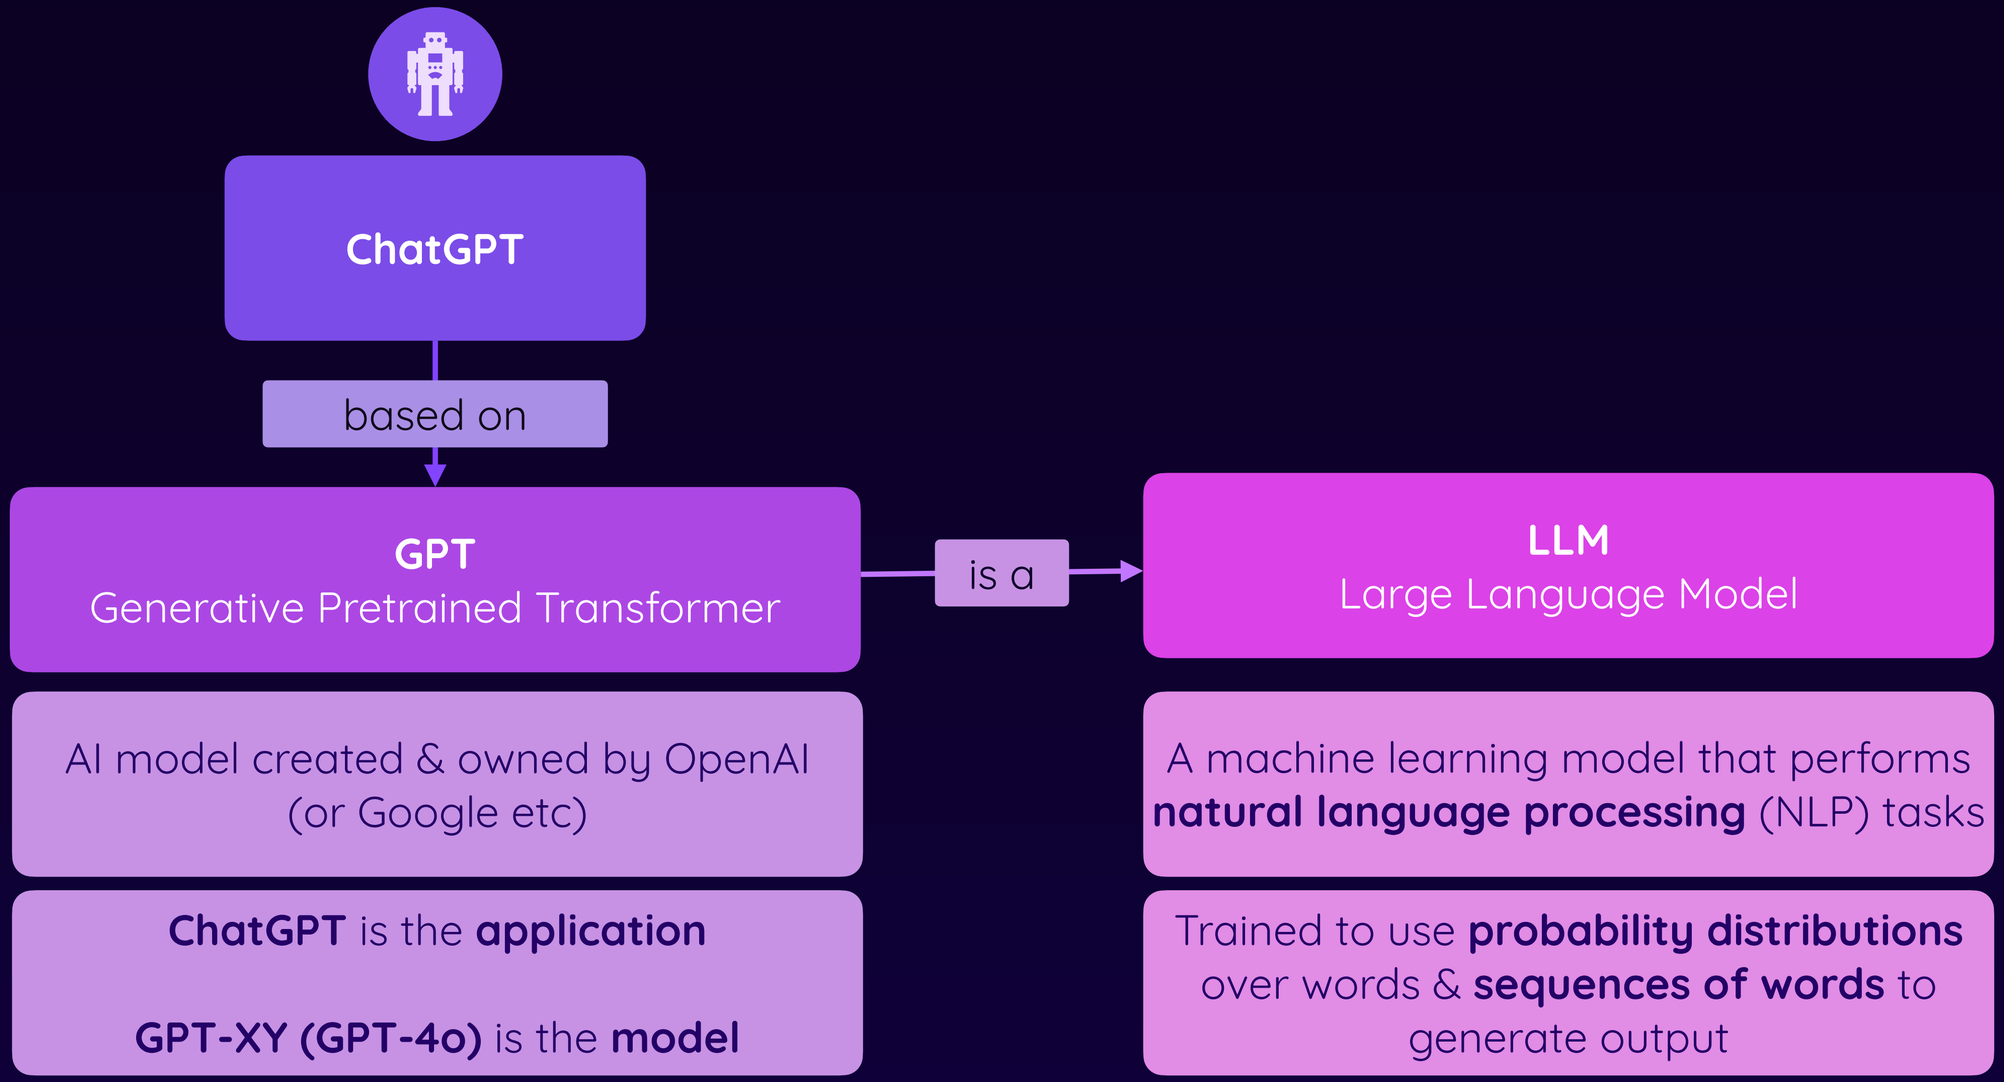

Model vs Application

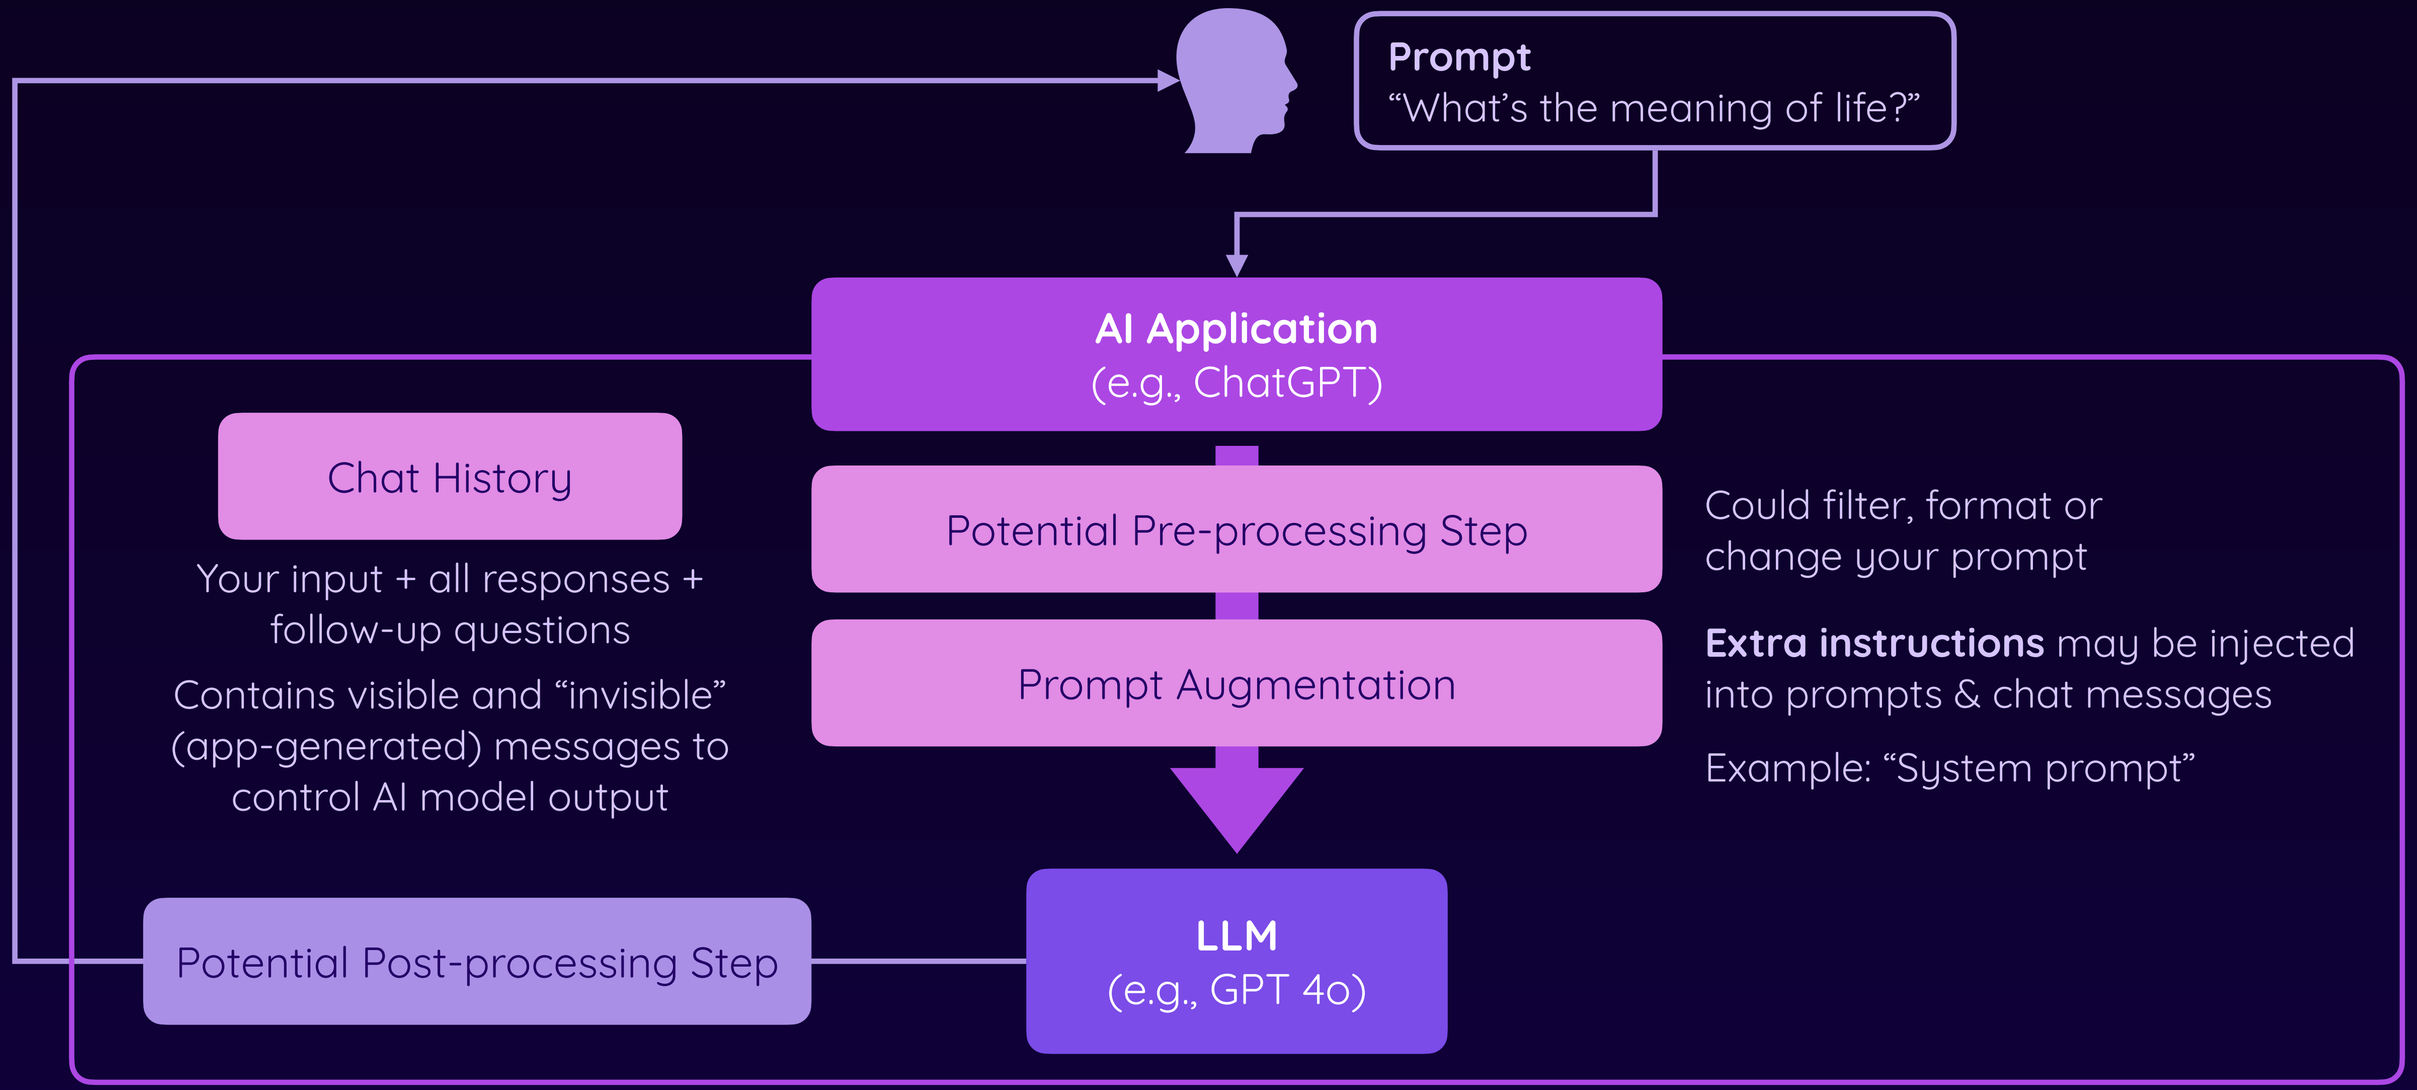

Augmented Prompts & Messages

AI applications (like ChatGPT) may **edit your prompt** or **insert extra (invisible) messages** into the chat history

Why?
To control the output of the underlying AI model

Examples

System Prompt

A special message that's injected into the chat history

Aims to define the general "behavior" of the AI model

Example: "You are a friendly and helpful assistant. If the user asks for news or recent developments, reply with 'I don't know that, sorry'"

Tools

Covered in-depth later!

Some AI applications expose "tools" to the AI models

The AI models may request the use of a provided tool to better fulfil a user request

Example: "You can search the web, if needed. Reply with WEB SEARCH: <search term> if you want to use this tool"

Retrieved Data (RAG)

Covered in-depth later!

For certain requests, personal or restricted data may be required

AI apps can inject retrieved data into a prompt to make it available

Example: Fetch and inject PDF document contents

Interacting with AI via APIs

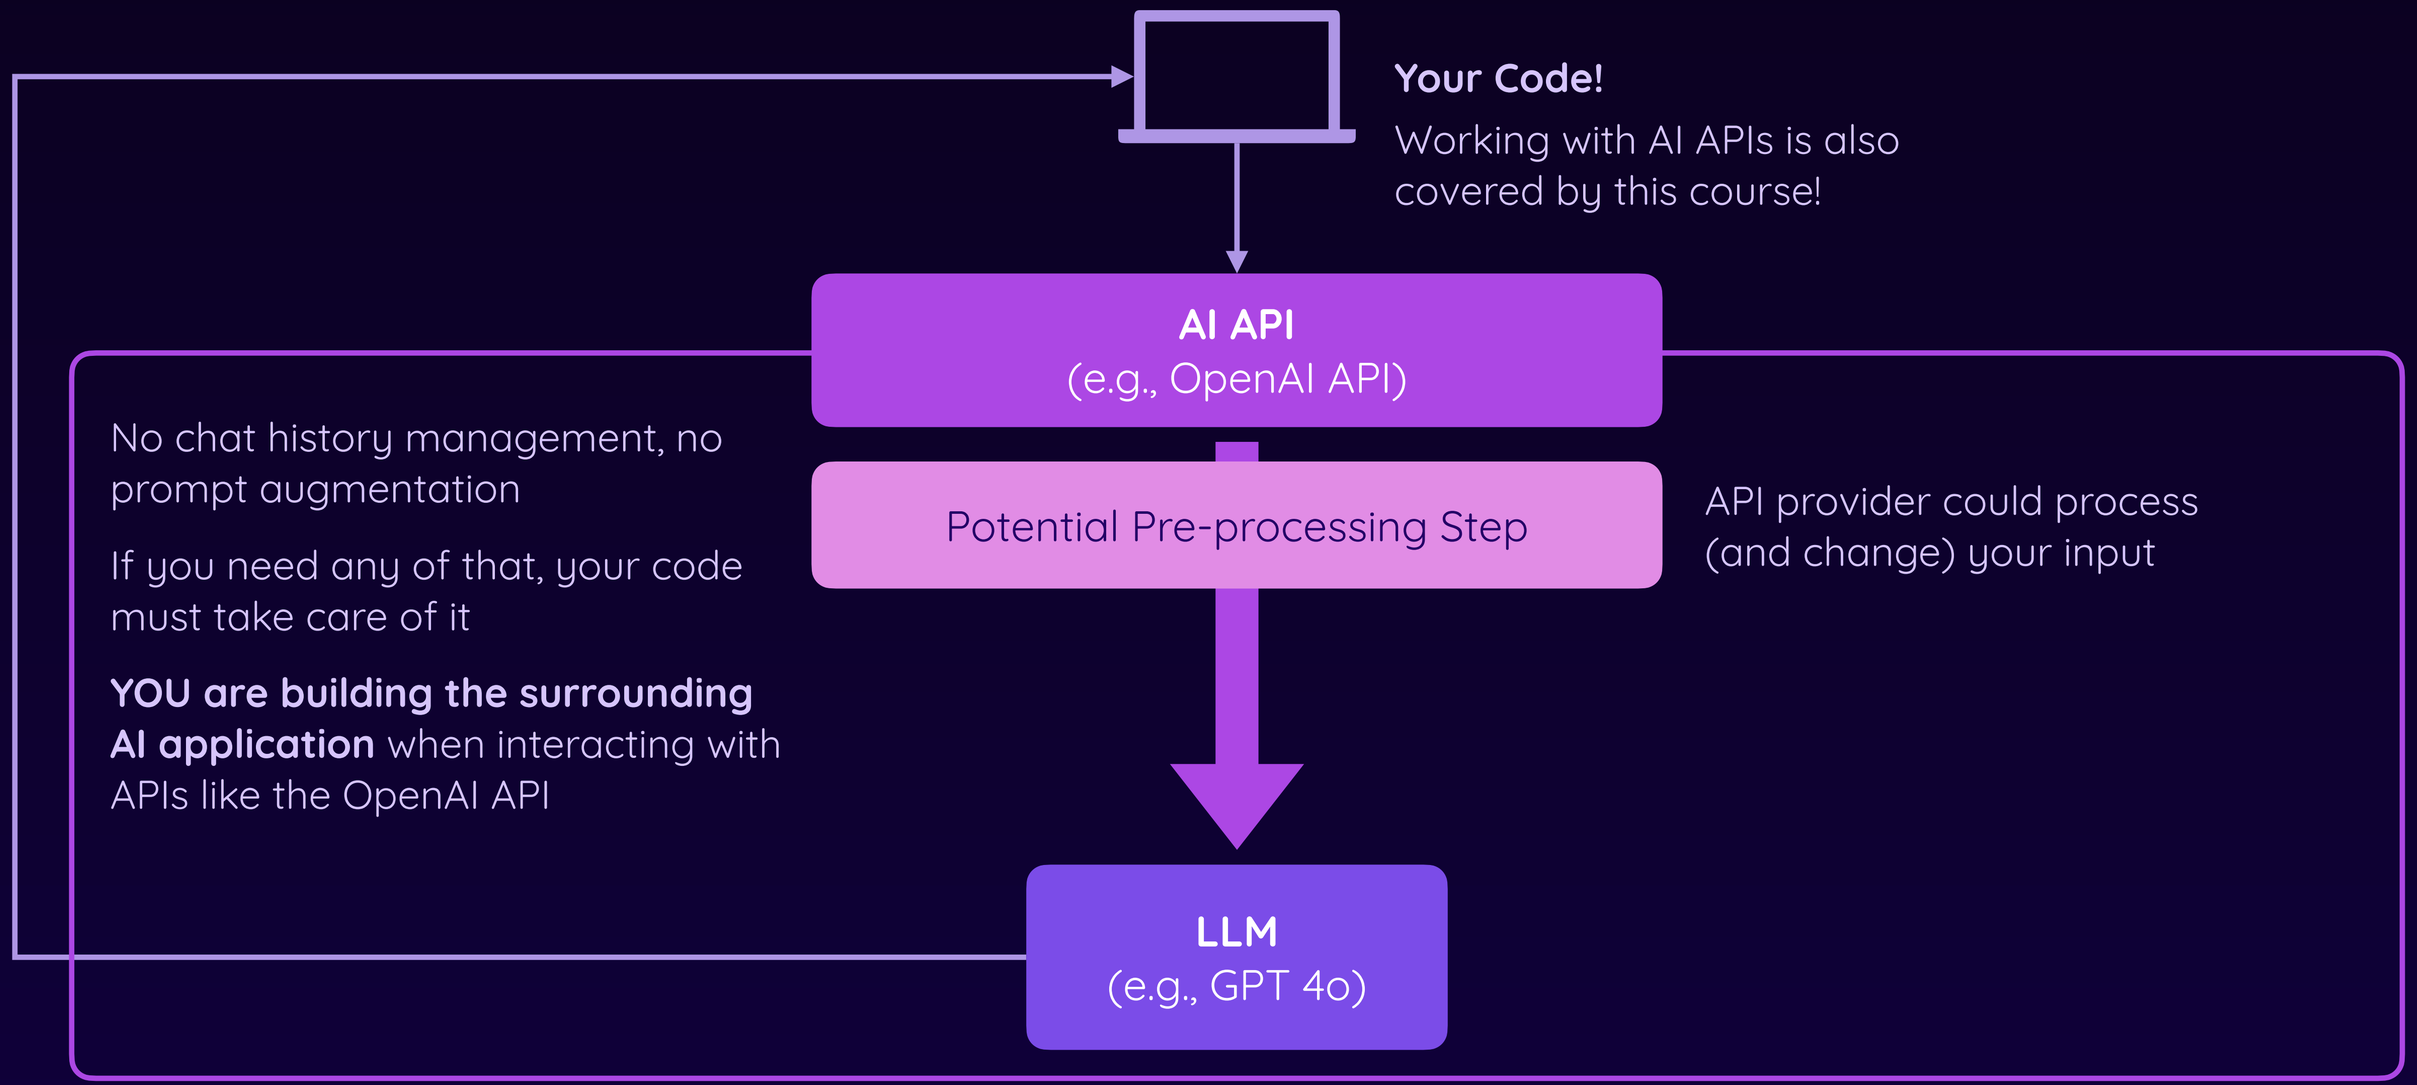

Model vs Application

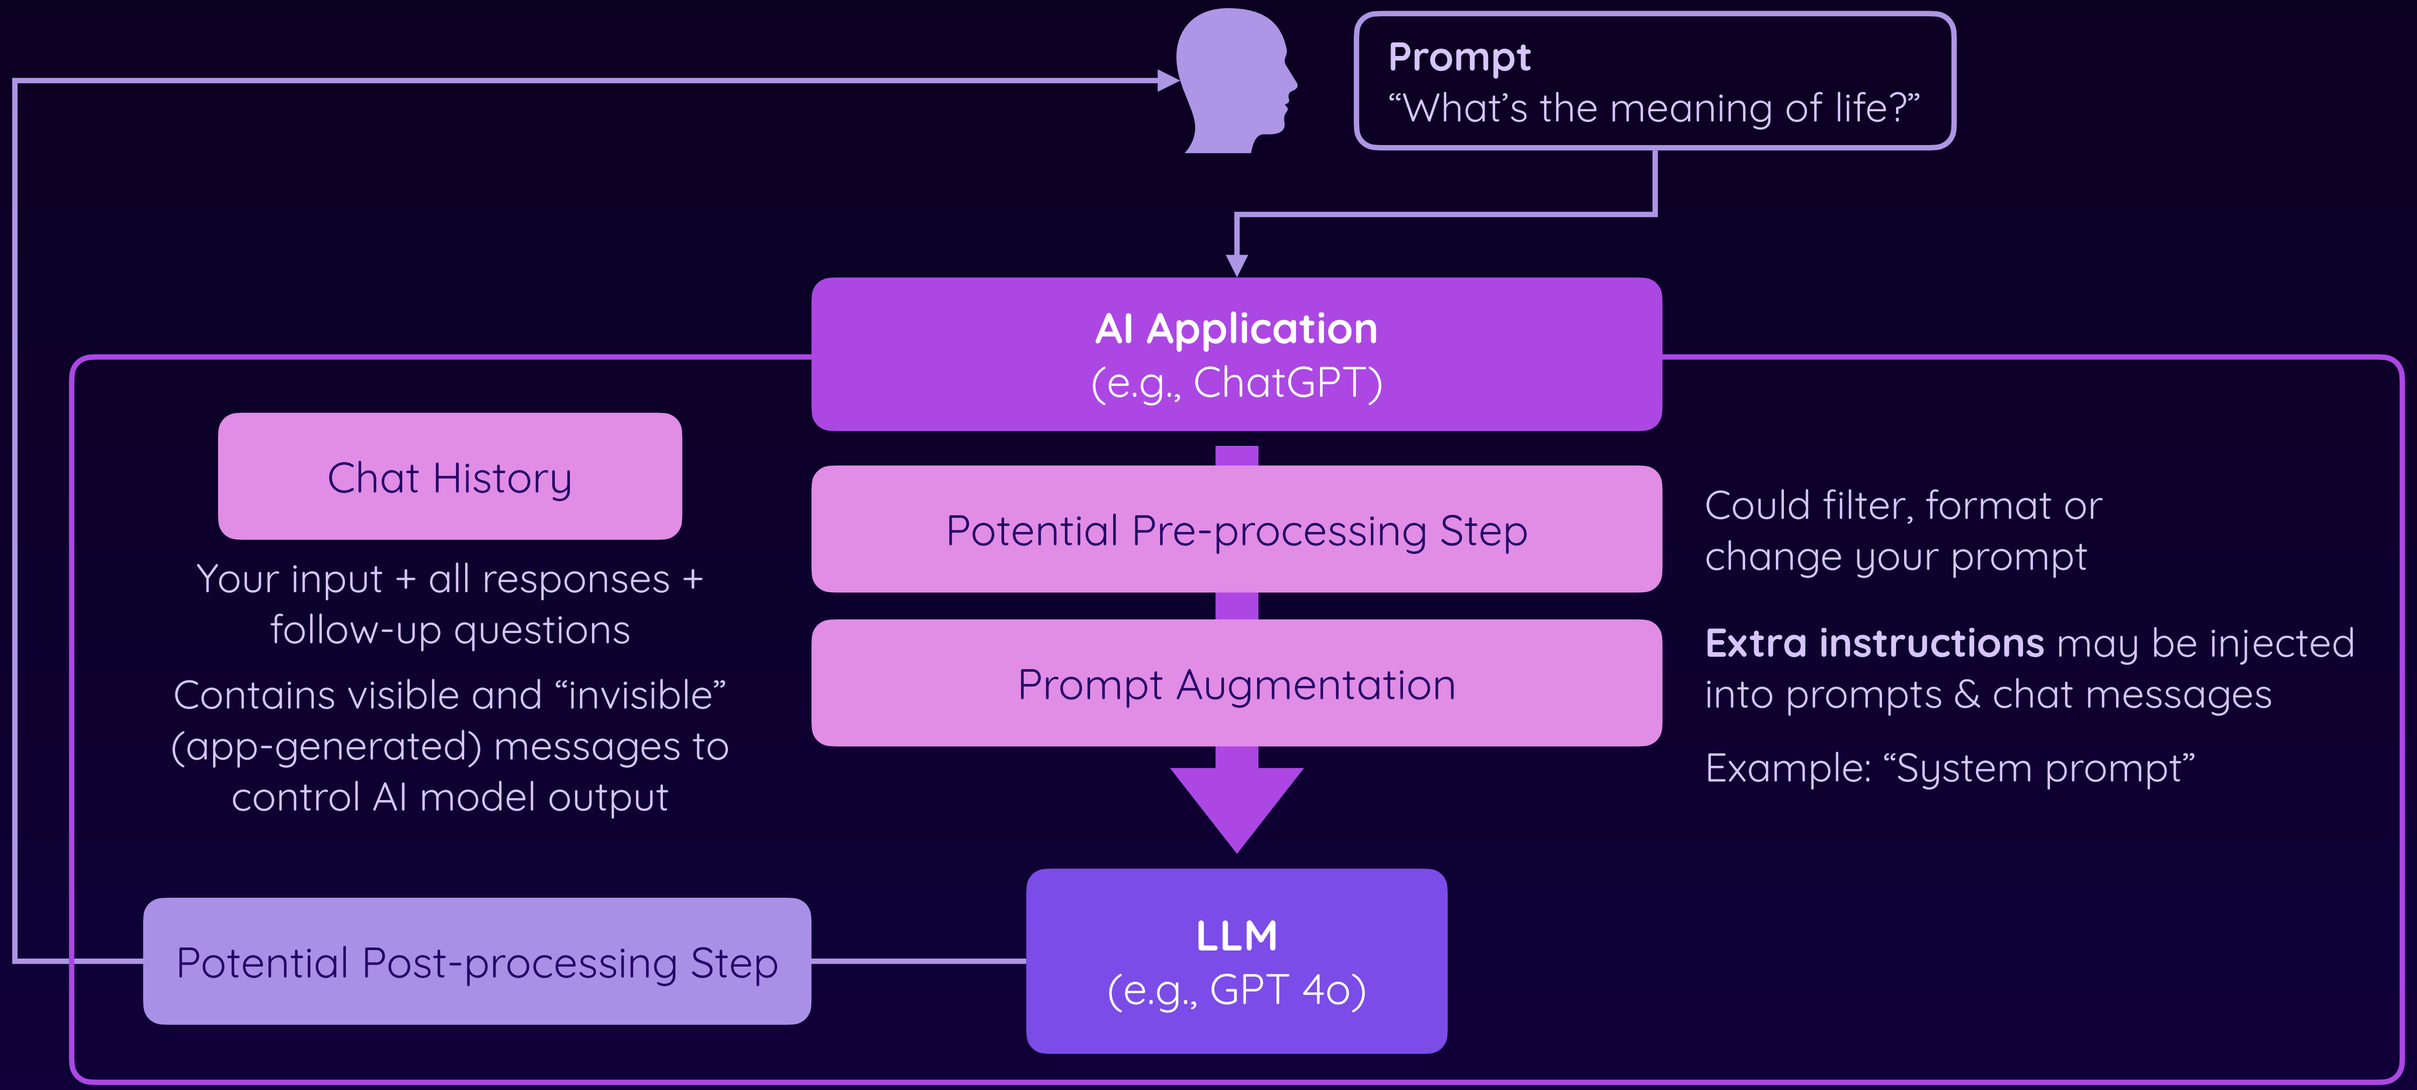

The Most Important Part: The Model

LLM
(e.g., GPT 4o)

The **Large Language Model** that generates the response text

LLMs Are (Huge!) Neural Networks

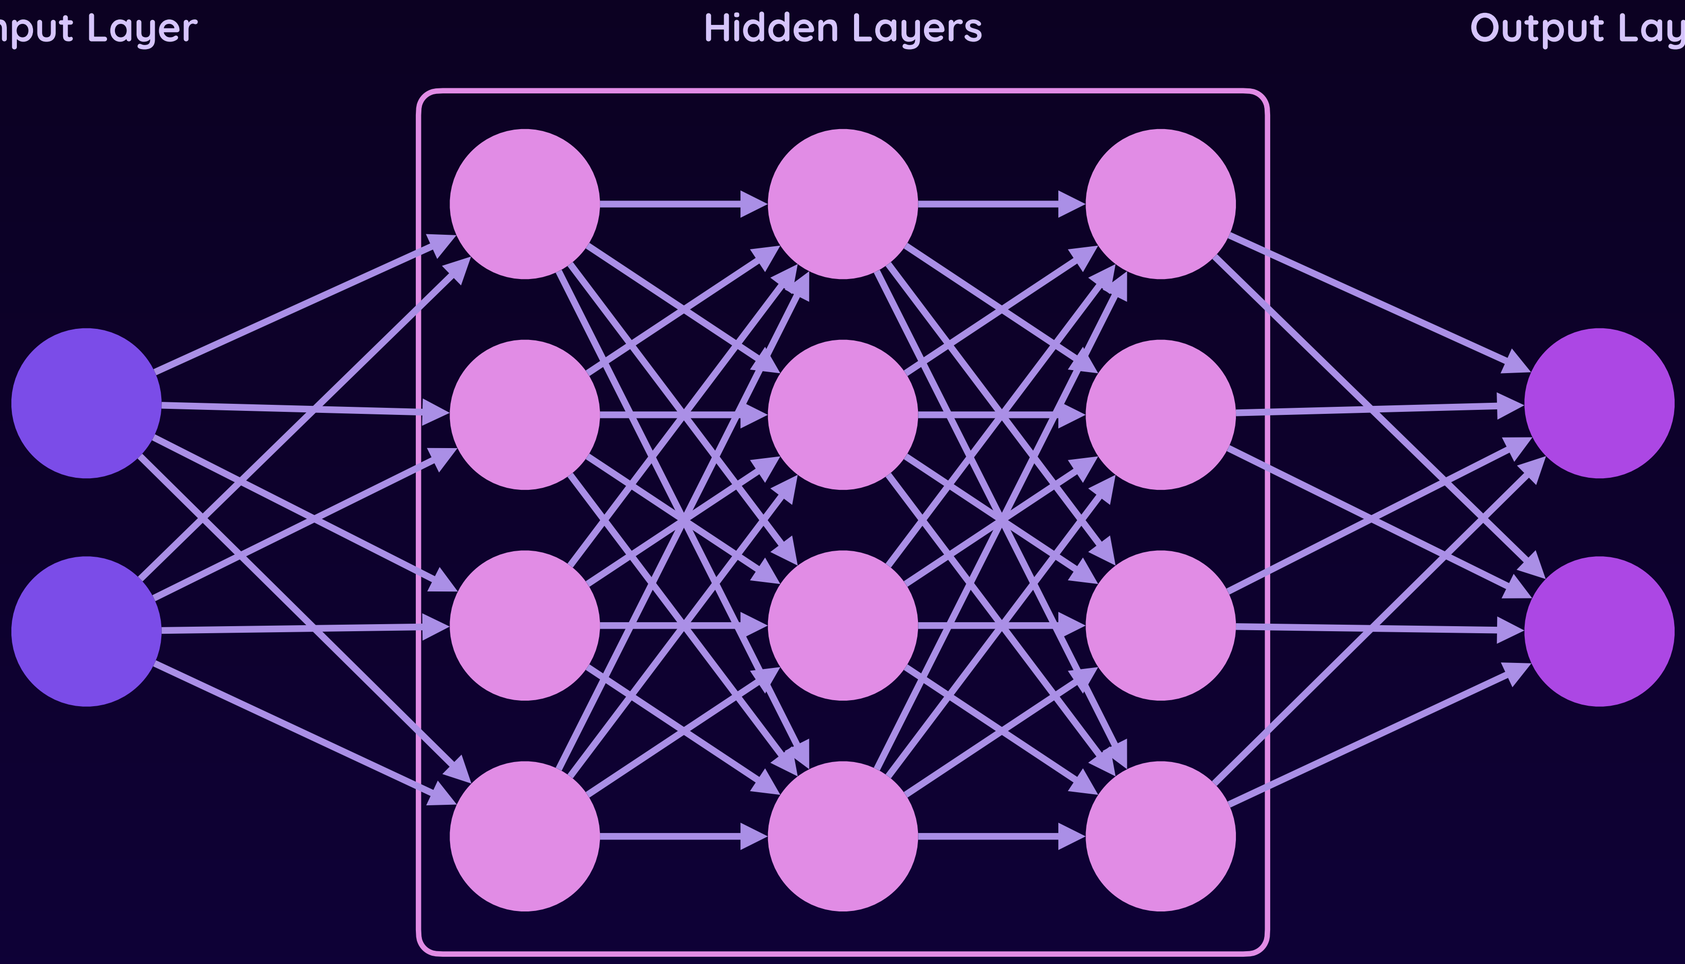

Each connection has a weight assigned to it

These weights are the so-called **parameters**—they are derived during the model training phase

Training LLMs Requires Text—Lots Of Text

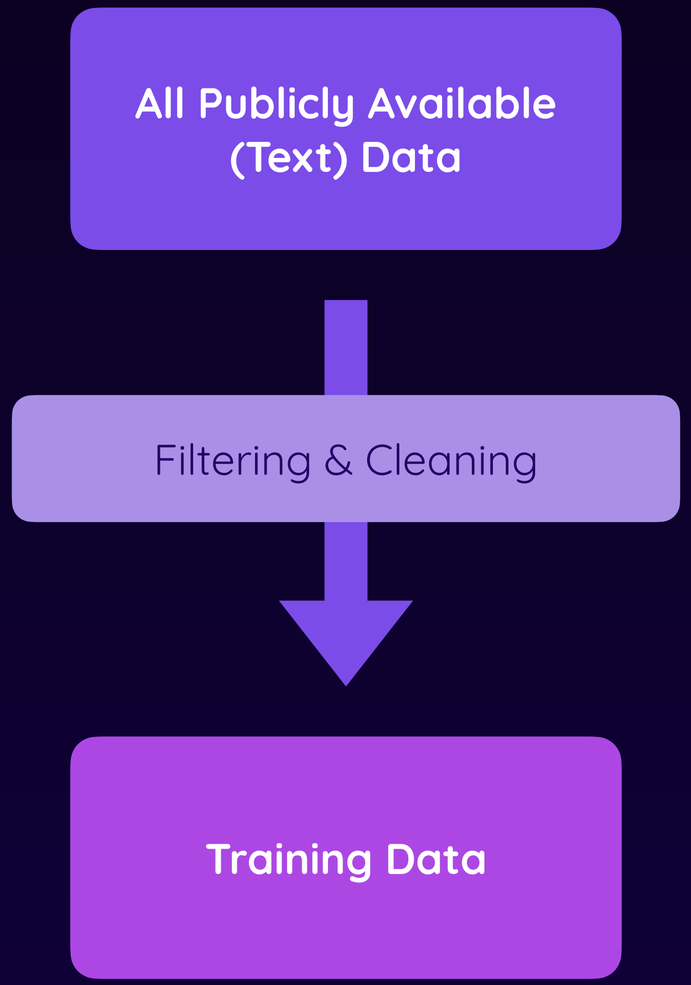

News articles, blog articles, forum discussions, ...

Example: Filter out adult content, illegal content, ...

Training a Foundation Model

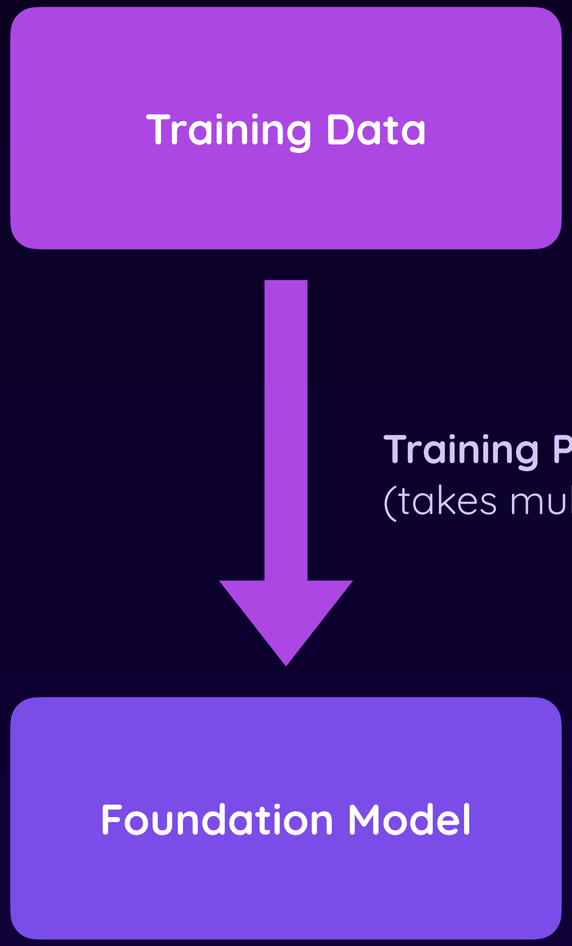

Training Process
(takes multiple months for large models!)

A LLM that “understands” the relation between words & sentences

This model is able to generate text based on a received input sequence of words

Understanding Relations Between Words

Large Language Models “understand” the relation between words

This allows them to predict future words & complete sequences

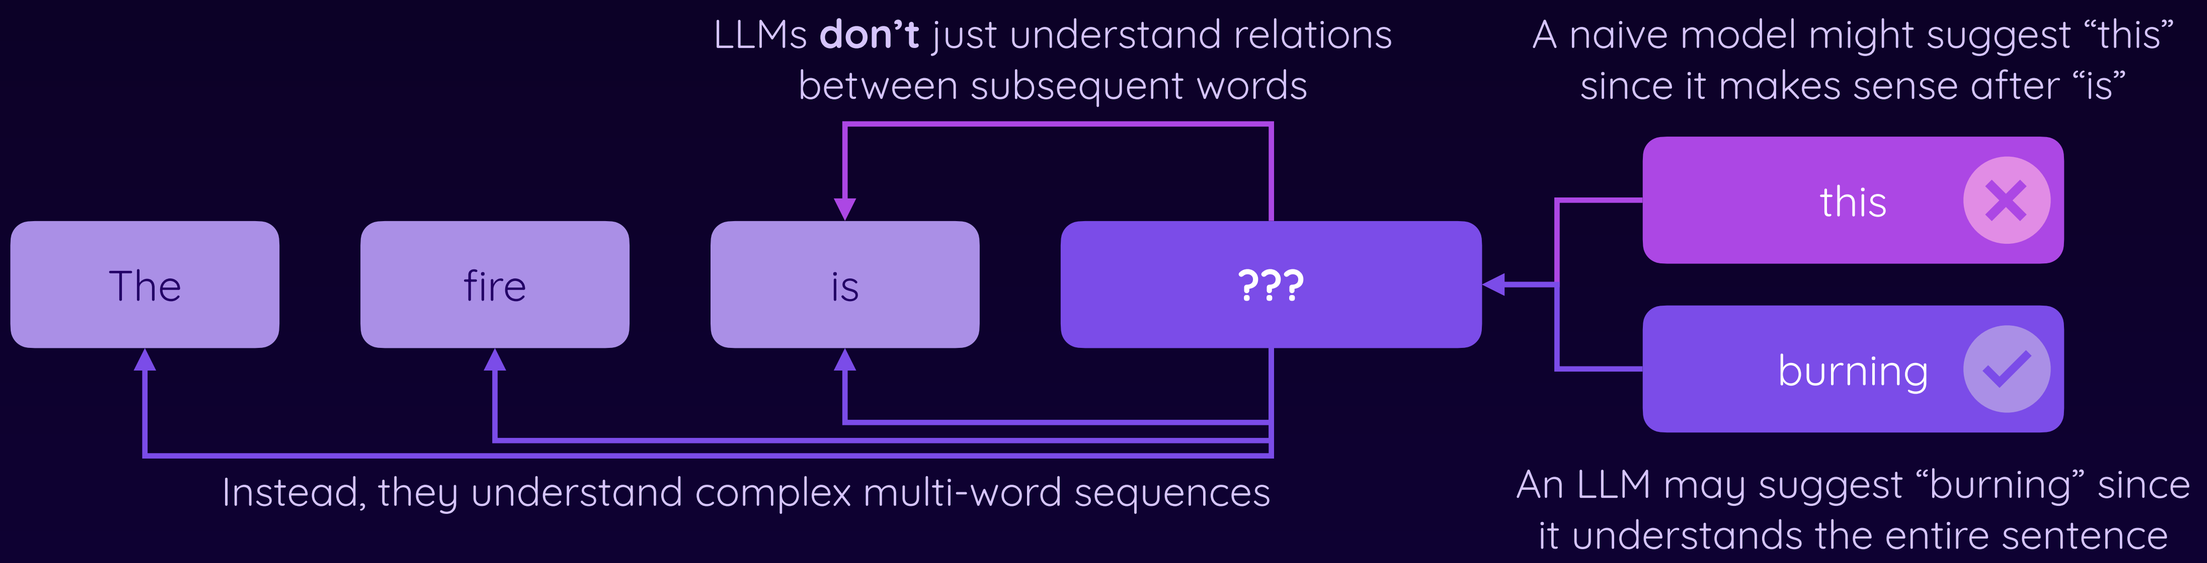

LLMs Are Trained With (Lots Of!) Text

LLMs are trained with large amounts of text

From Text To Tokens

LLMs are trained with large amounts of text

LLMs Operate on Vector Embeddings

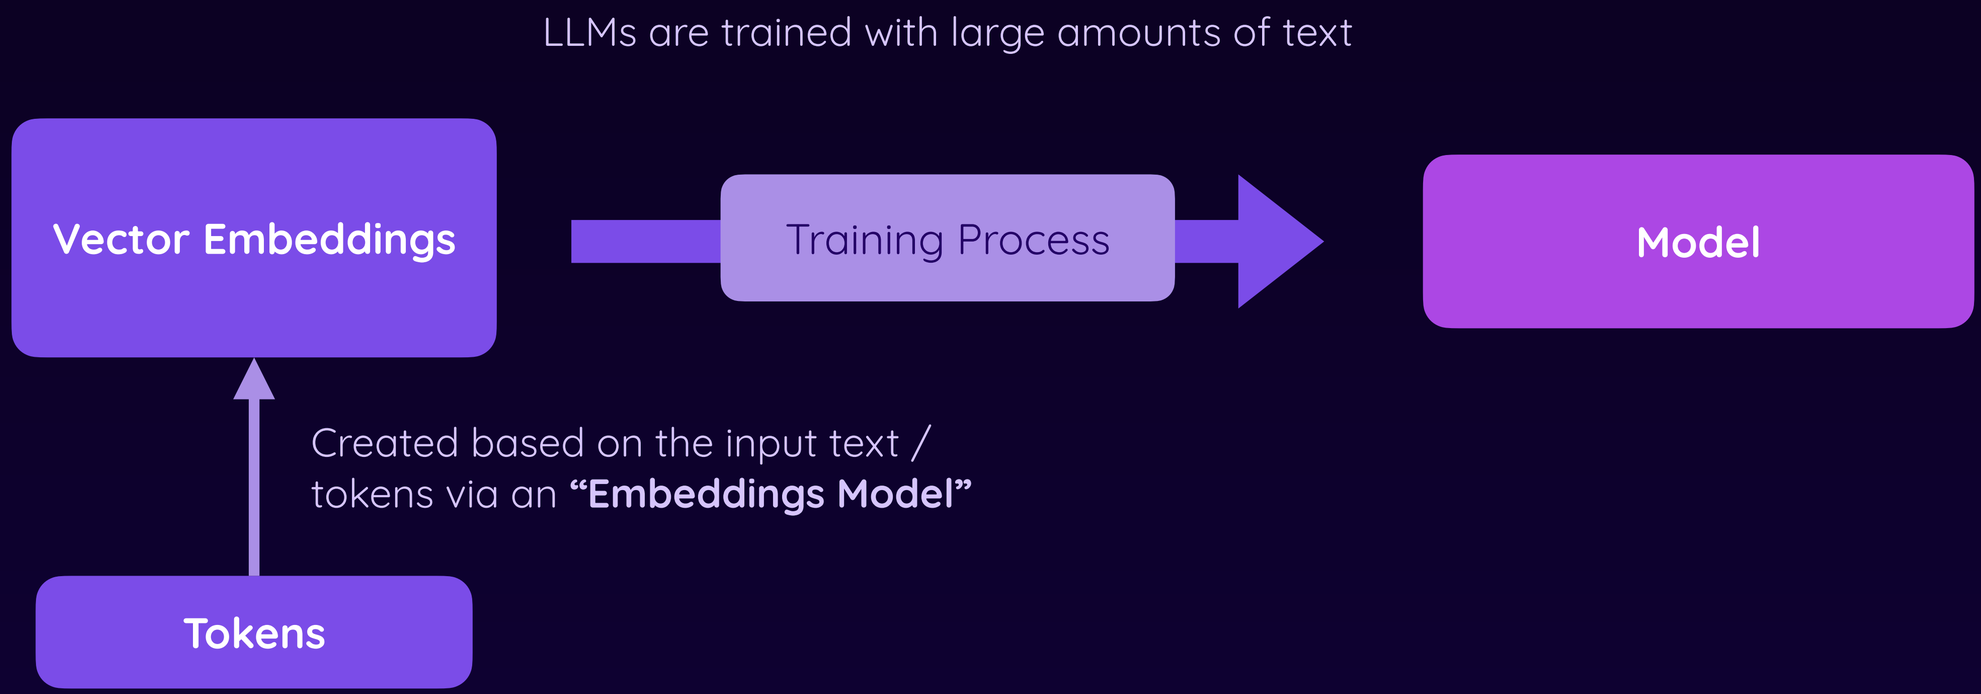

Vector Embeddings Represent Relations

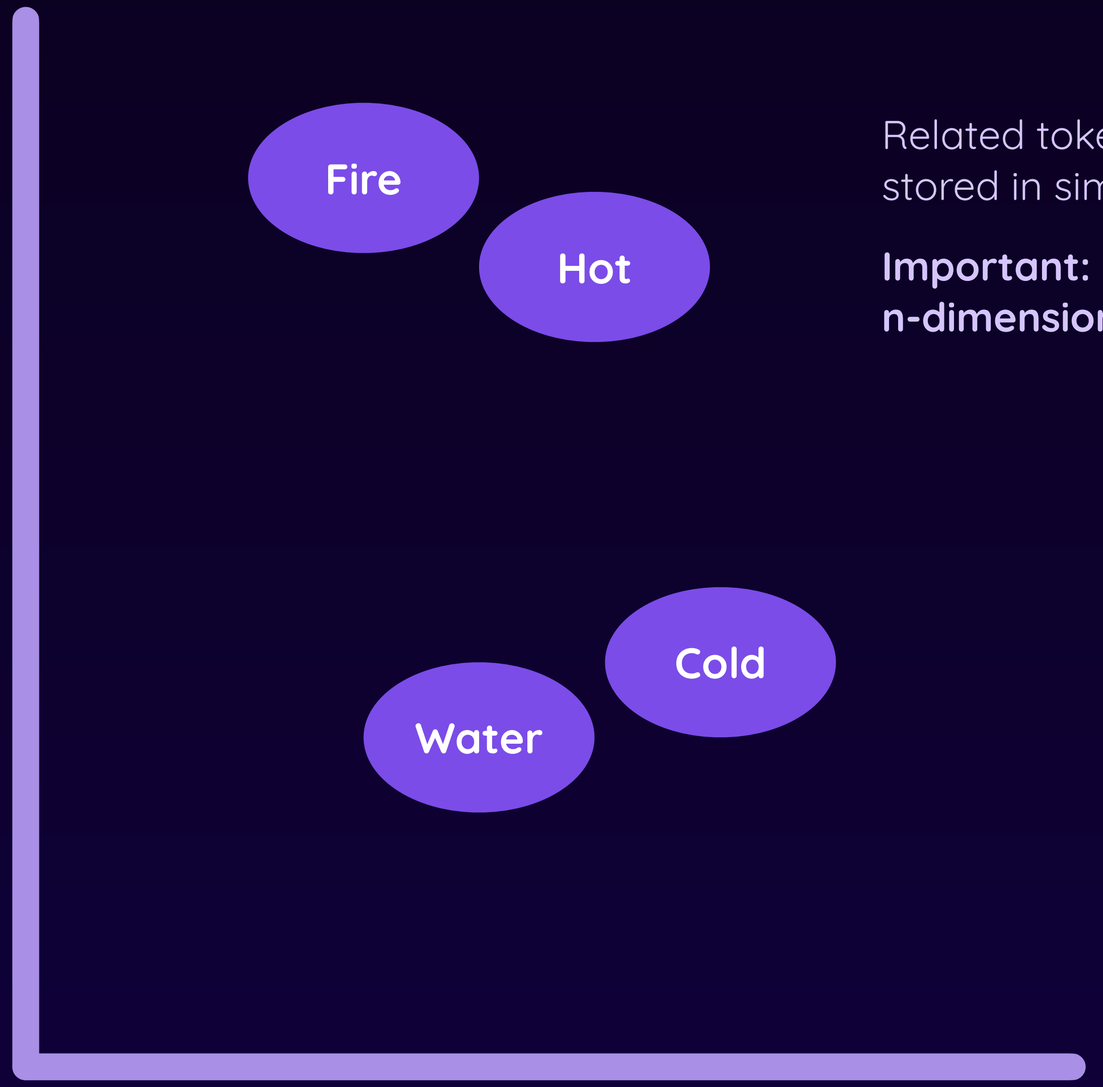

Related tokens / words are stored in similar places

Important: In reality, it's a n-dimensional space!

LLMs Generate Tokens

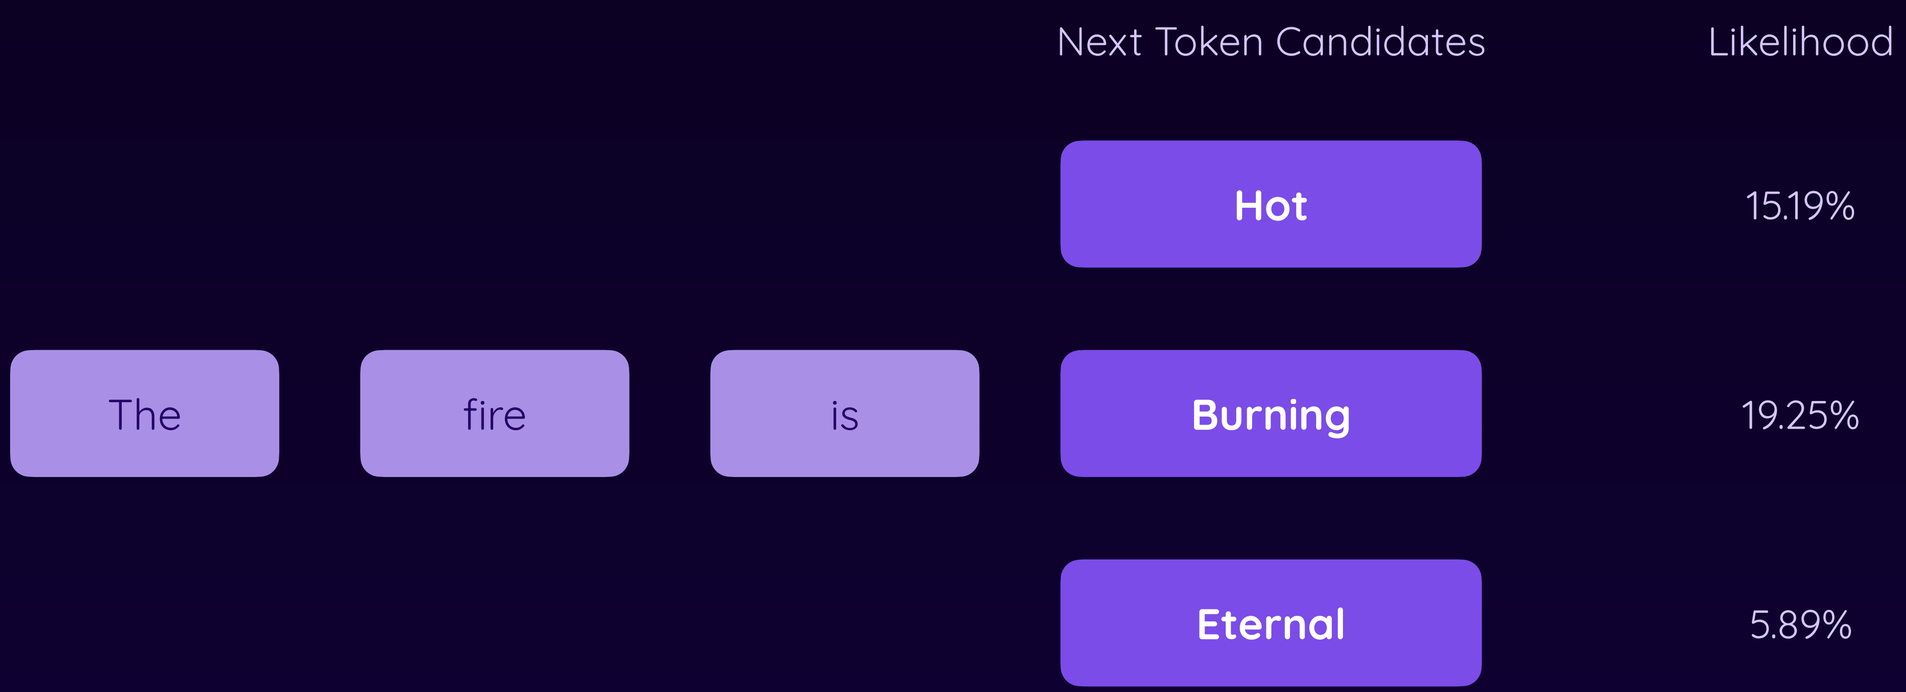

The application then chooses one of those candidates -
based on the probability and possibly also other settings

Foundation Models

After finishing training, you have a **foundation model**

Foundation Models are LLMs that can complete word / token sequences

They don't necessarily make for great AI assistants or chatbots, though!

User	What's the meaning of life?
Foundation Model	That's the question. And also: Why do we even exist? Will we ever find out?

From a helpful **AI assistant**, you **would expect an answer, not just a completion.**

From Foundation Models To Assistants

This phase is called **“pre-training”**

The result of the pre-training phase is a **foundation model**

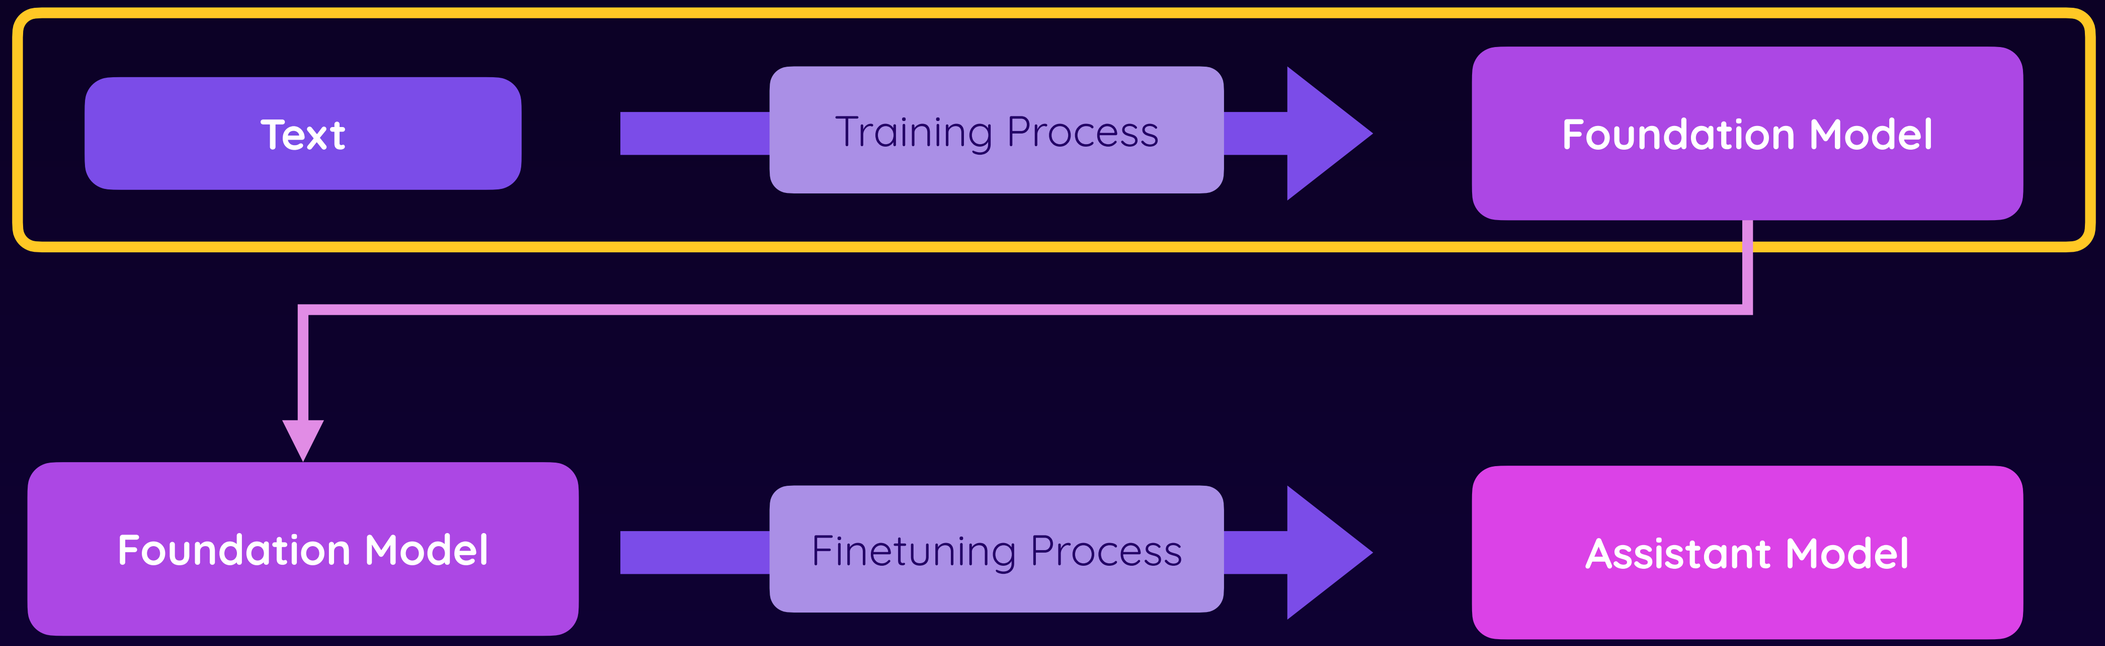

During finetuning, the foundation model is used as a base to train a **finetuned assistant model**

Finetuning is performed by adding (human-created) training data that simulate human<=>AI assistant interactions

Assistant Models

After finishing finetuning, you end up with an assistant model

Assistant models are still LLMs that predict future tokens

But they are finetuned to generate tokens in-line with the finetuning data

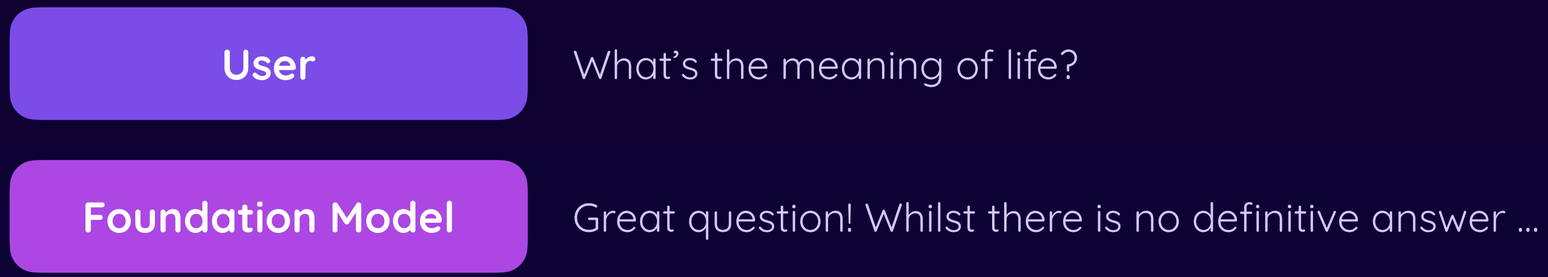

When you interact with ChatGPT etc., you're interacting with assistant models
(also called "instruct models")

Onwards To Reasoning Models

Some modern LLMs can “think” before they answer

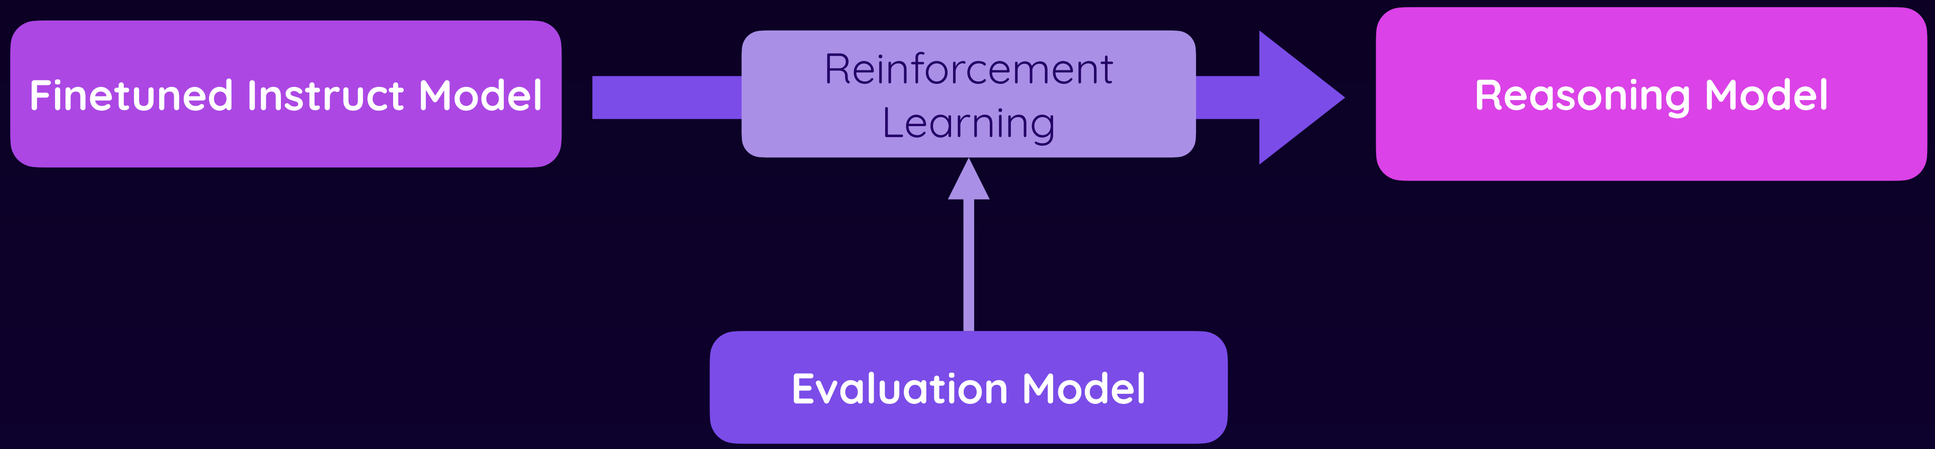

During the training phase, problems are given to the instruct model

The responses / results are than evaluated & graded by a separate evaluation model

That “feedback” is then used to adjust the LLM parameters until satisfying results are produced consistently

Going through an initial “thinking process” was “learned” during this training phase

LLMs Have Context Windows

Chat History

Your messages	Your attachments	Model Responses	Invisible Messages
Your questions & follow-up messages	PDF documents, images, ...	AI-generated responses	System prompts, other app-controlled messages

Context Window

This entire history must be processed by the LLM for every (!) new response it generates

It's the "available context" the model takes into account for its response generation

Context window: Each LLM has a **maximum amount of tokens** it can consider

Knowledge Cutoff

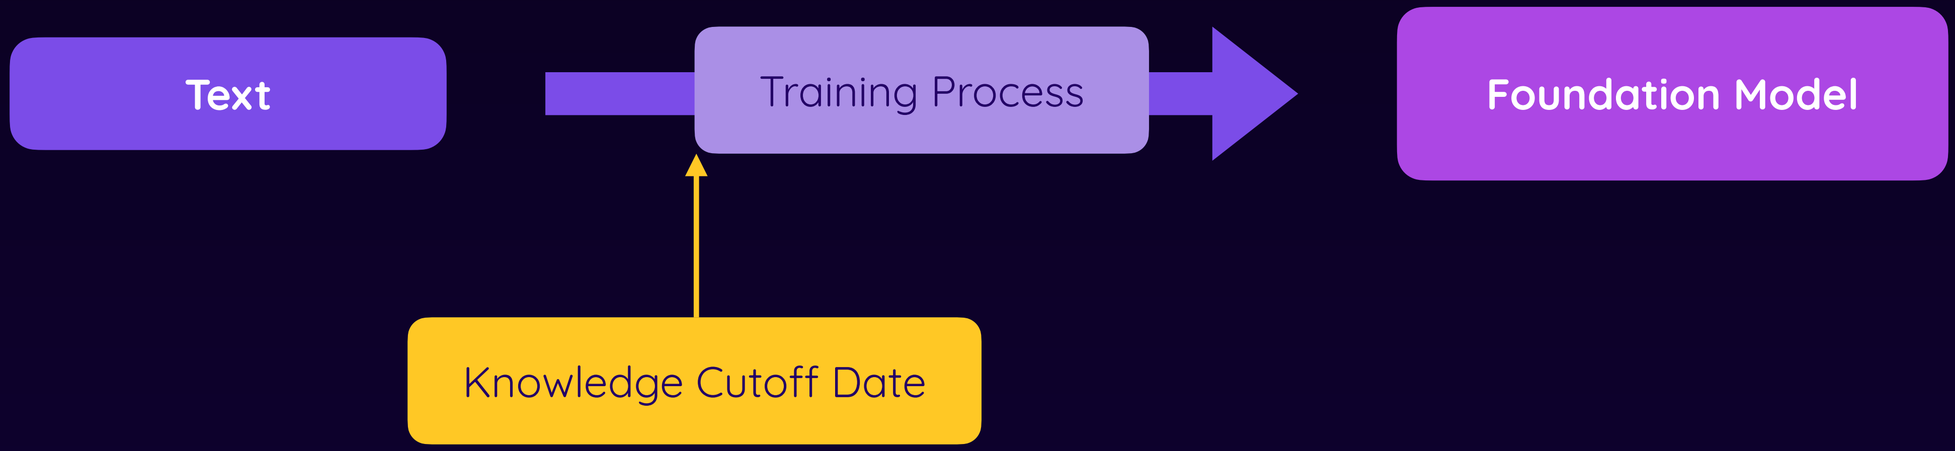

The model can't "know" any facts after the start of its training process

Solution 1: Prevent Hallucinations

Finetune data to "detect" questions it's likely not able to answer

Generate a generic response ("Sorry, I don't know that")

Solution 2: Expand Knowledge

Give AI application access to web search & finetune to detect relevant questions

Application can perform web search & include results in context

How AI Assistants Use Tools

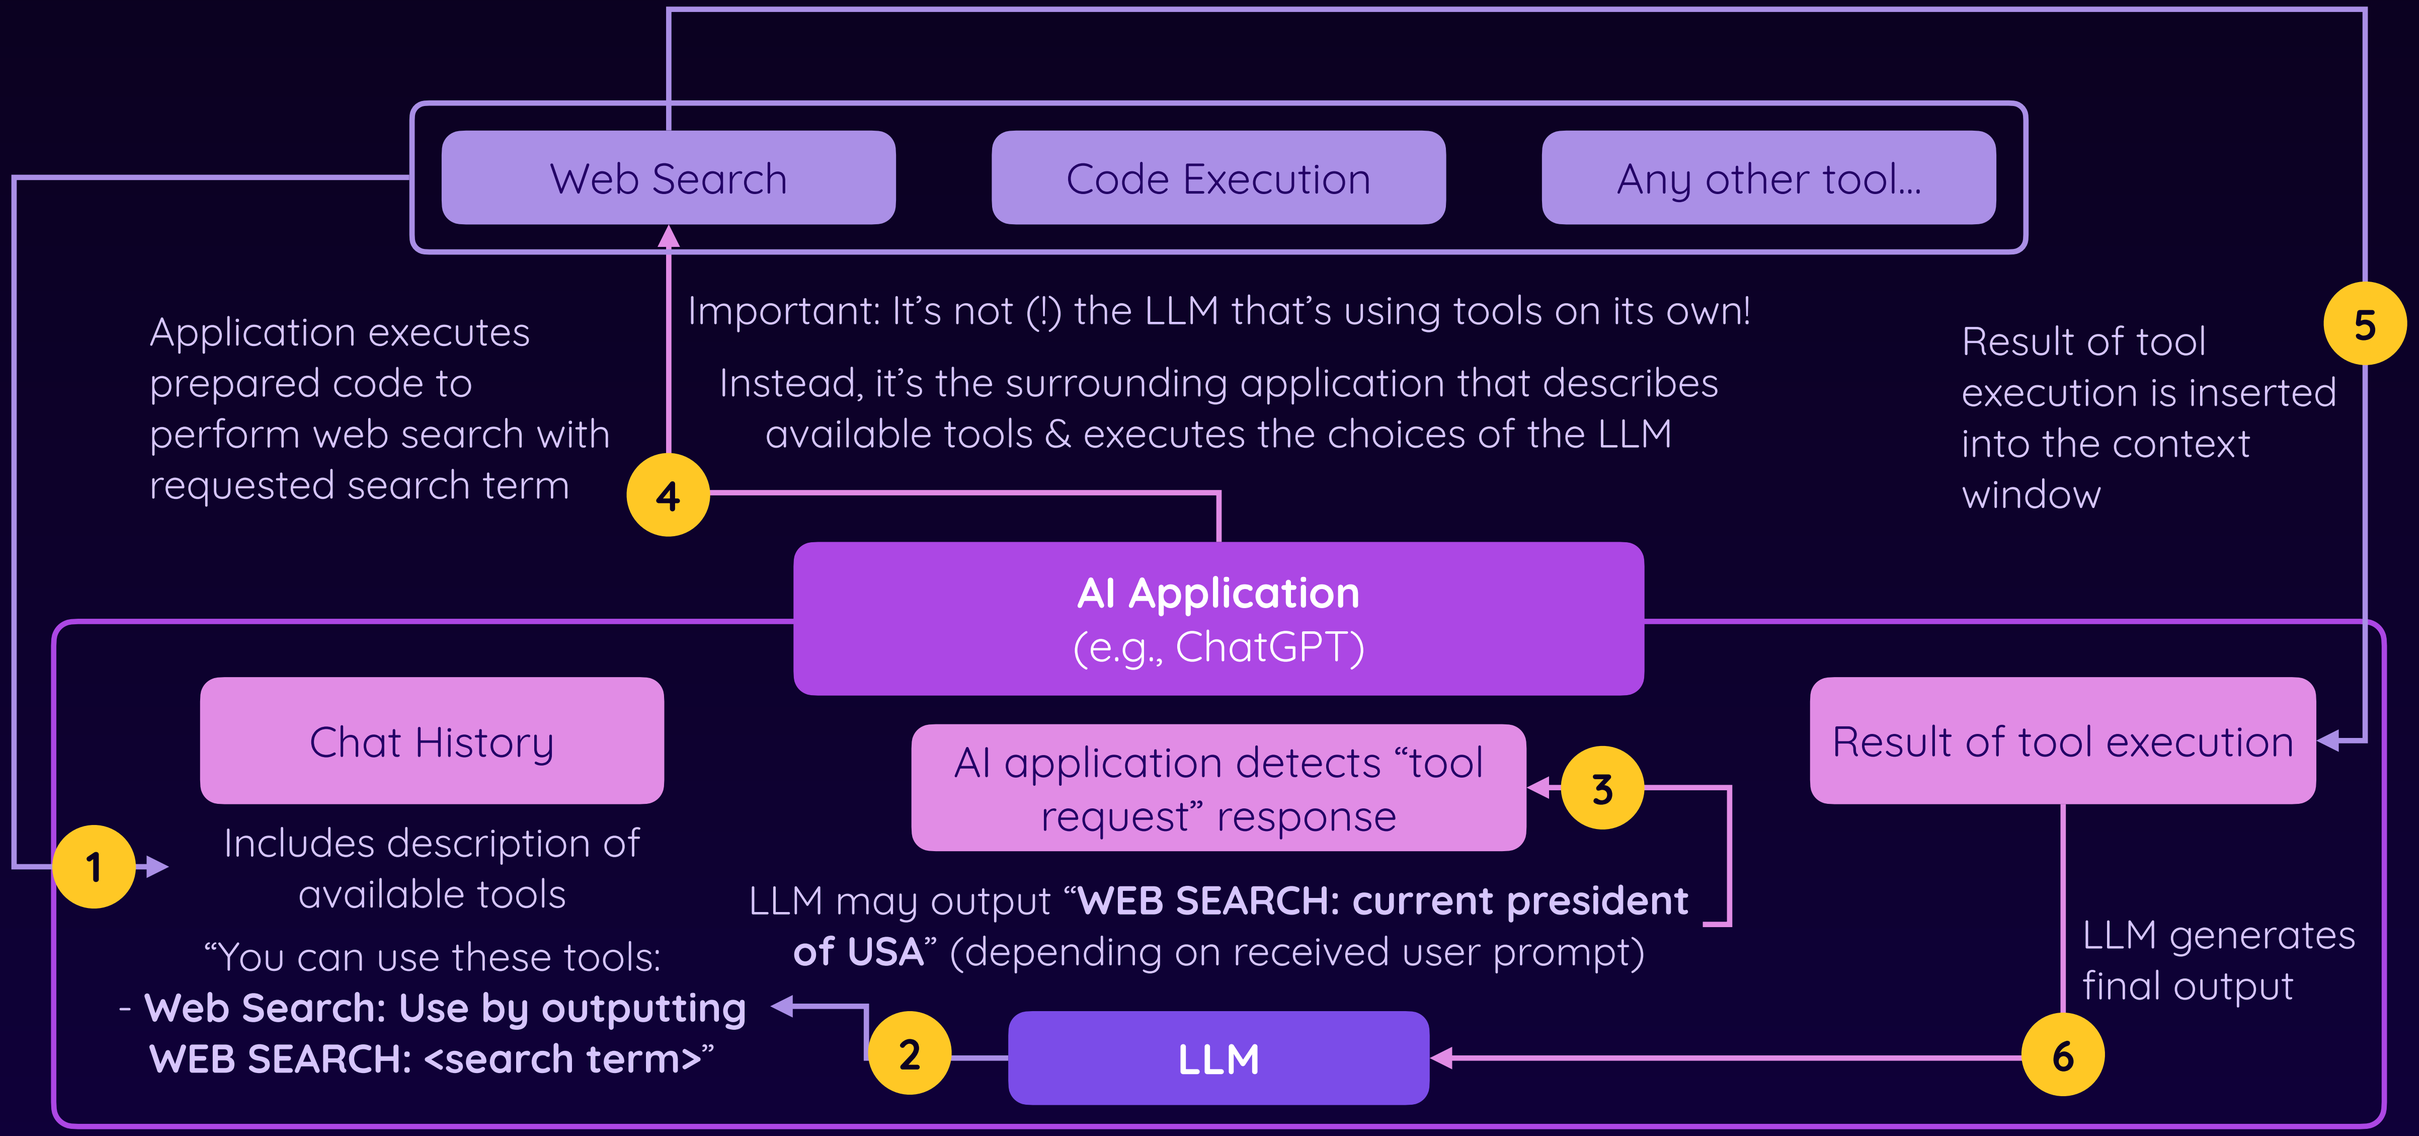

It's All Highly Dynamic!

ChatGPT

Still an AI chatbot but ...

... more models

... more capabilities

... older models are gone

AI chatbots & their capabilities keep evolving

And different AI model providers offer different features

ACADEMIND

Running LLMs Locally

ChatGPT & Co vs Self-hosted Solutions

- ChatGPT & Co Disadvantages
- Exploring Open Models
- Understanding Quantization
- Running LLMs Locally

Typically, the weights / parameters are “open”

Not the code / algorithm

ChatGPT & Co Disadvantages

Cost	Availability	Privacy
Most models are not accessible for free	No access without a (stable) internet connection	You're sharing your data with the model providers
Having multiple subscriptions can add up	Services may go offline (e.g., due to demand)	Usage restrictions may apply
API access is charged based on usage	Usage limits may apply	

But: If you want / need to use the **most powerful, versatile and knowledgeable models**, there's (almost) no way around ChatGPT & Co

Almost: There is DeepSeek - an open-source model provider. But running / **self-hosting large LLMs is very difficult & costly**

Onwards To Open-Source Models

Most open-source models are **smaller** (i.e., less parameters) than the most capable proprietary models

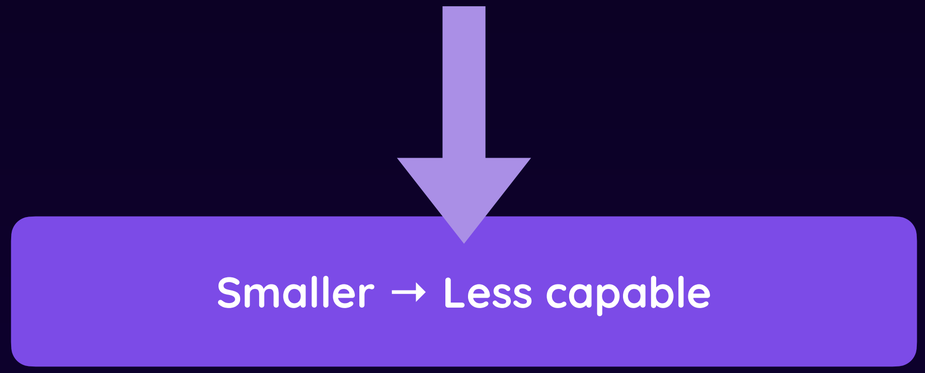

Depends on your use-case!

For many tasks, a small, locally running model may be equally good or even better than ChatGPT & Co

- Summarization Tasks
- Data Extraction Tasks
- Local Knowledge Tasks

Open-Source Models—An Overview

Whilst most OpenAI models are proprietary (i.e., you can't self-host them), there are **plenty of popular & powerful open-source models** available

Meta's Llama Models

Google's Gemma Models

DeepSeek's Models

And many others!

Huggingface is a great place to explore all those open-source models in detail!

Understanding Quantization

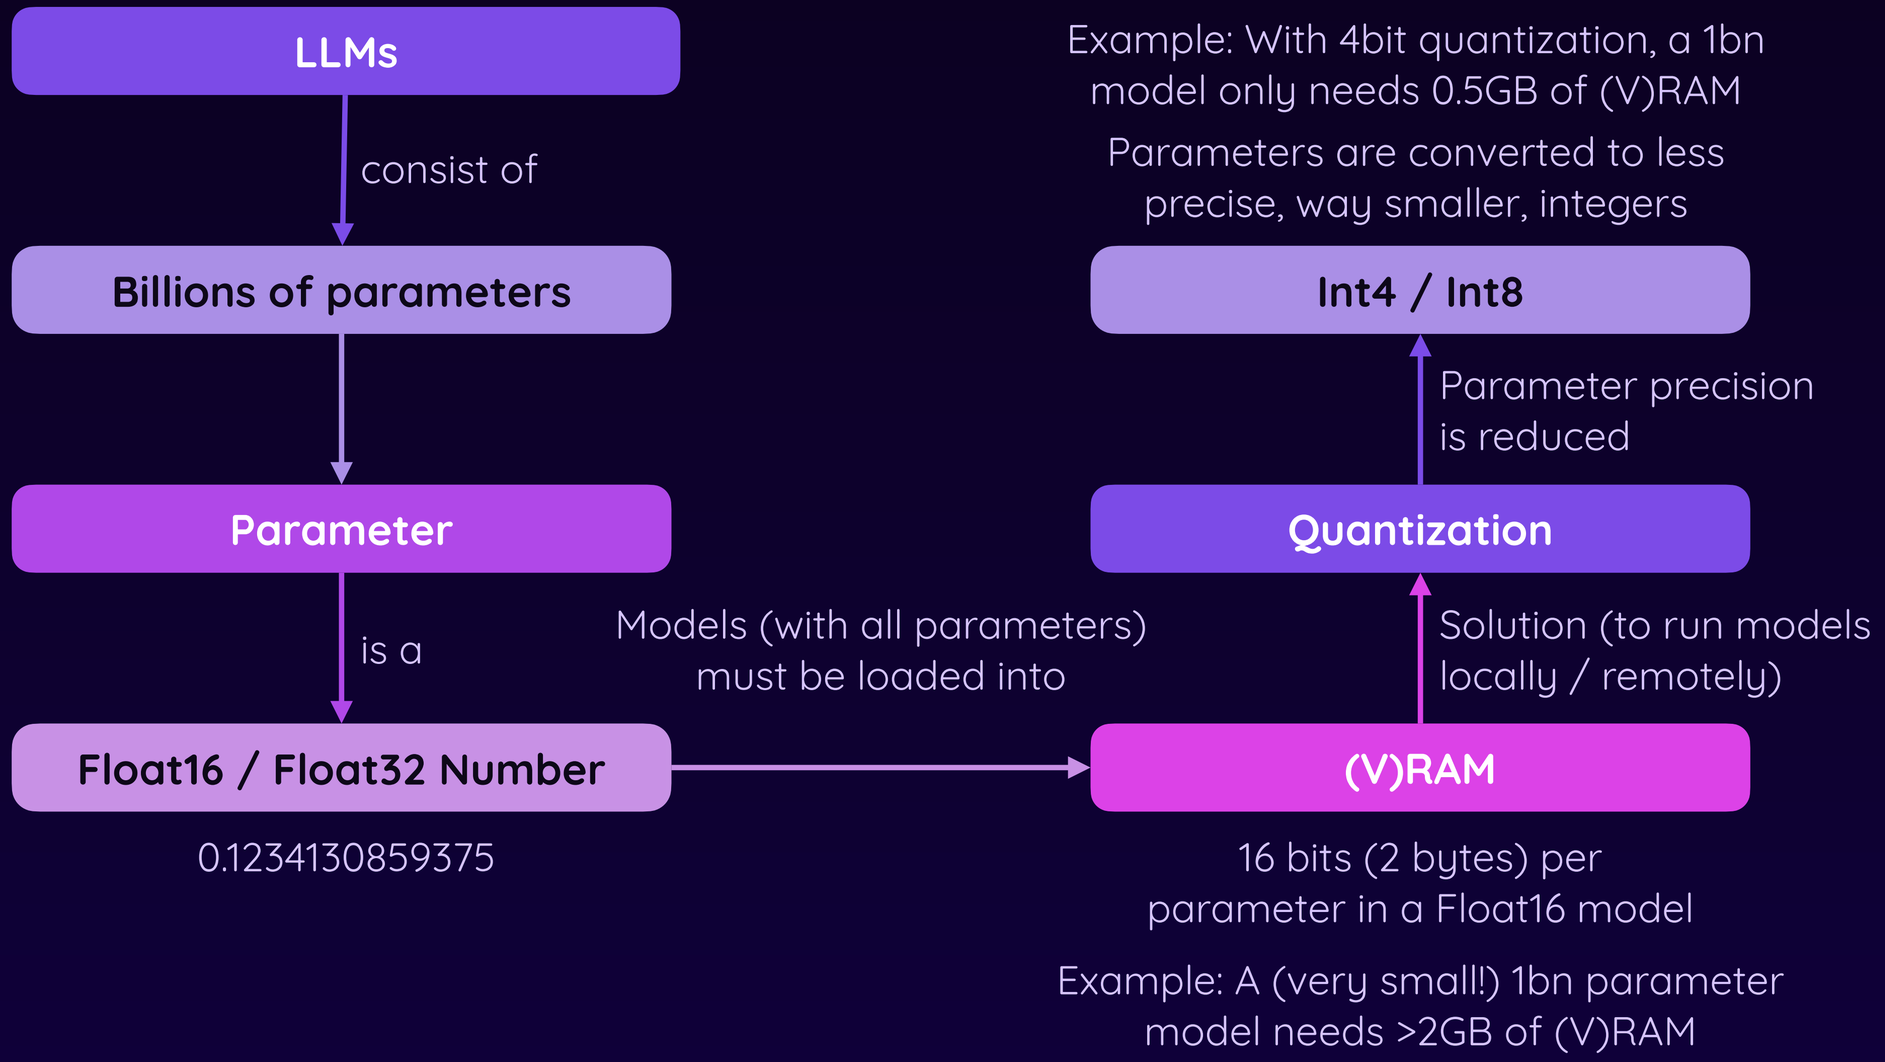

Running LLMs Locally

There are different solutions that make running open-source LLMs locally a breeze

LM Studio

Ollama

llama.cpp

You can also use open-source models programmatically via Ollama & Co APIs OR via Huggingface transformers library

Self-hosting LLMs

You can also run (host) open-source models on owned / rented servers

On-premise / VPS	Managed Service
You own / rent & configure the server	You use a managed service like Amazon Bedrock
You manage the software and install Ollama / LM Studio etc.	You select models & features → No setup required
You configure the network	No technical expertise required
You only pay for the server	You pay for the usage

Important

Self-hosting LLMs on a remote server is not cheap and not necessarily trivial

You will need at least somewhat capable hardware and go through various setup steps to run and expose LLMs from a remote server

ACADE MIND

Prompt Engineering

Ensuring Good Results

- What & Why?
- Understanding Key Techniques
- Best Practices & Recommendations

What Is Prompt Engineering?

What Is A Prompt?

What Is Prompt Engineering?

The Skill Of Crafting Great Prompts

Why Prompt Engineering?

Good Prompts

=

Good Results

Main Prompt Engineering Goals

Control Output **Content**

A LinkedIn Post

+

Control Output **Format**

Plain Text, Markdown, JSON, ...

Main Prompt Engineering Goals

Control Output **Content**

You want the LLM to generate content that you can use with as little modifications as possible

Of course, you also want a response that's correct, contains no hallucinations and meets any other requirements you might have

+

Control Output **Format**

For some (but not all) use-cases, you also might care about the format

You might want a response formatted as JSON or markdown

Or maybe you need text that's structured as a list of bullet points

What Defines A Good Prompt?

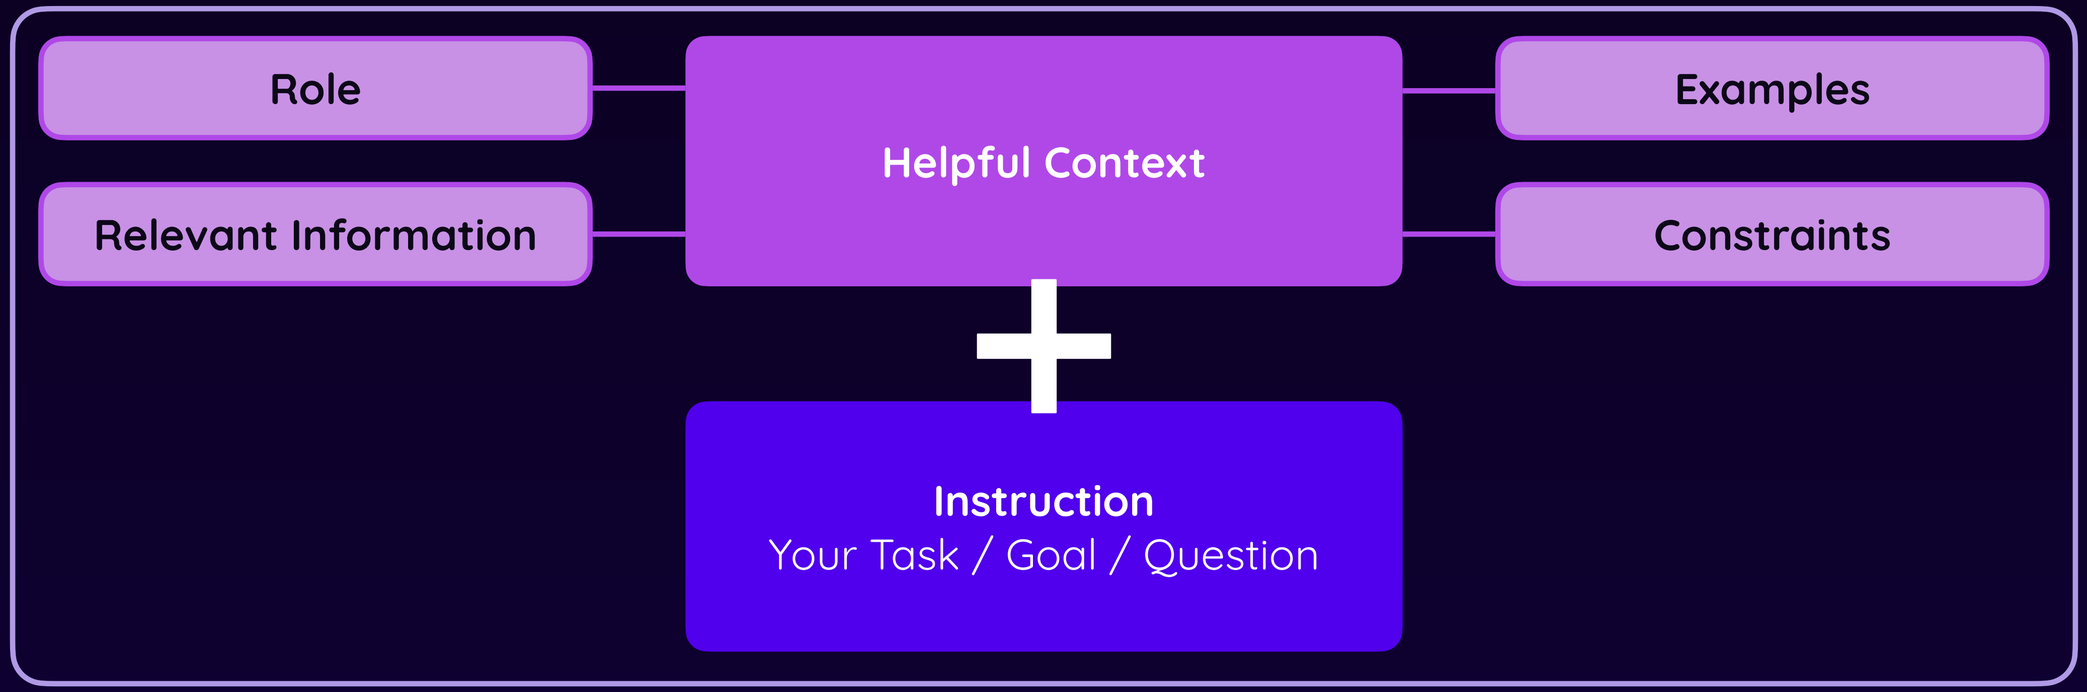

A good prompt includes a detailed task description and helpful context

Irrelevant information must be avoided

Complex (multi-task) prompts should be avoided

Keep In Mind

ChatGPT Output Is A Starting Point

Fine-tune & adjust as needed

Adding Meaningful Context

Prefer **short**, focused **sentences**

Add important **keywords** & avoid unnecessary information and ambiguity

Define the **target audience**

Control **tone, style & length** of the output

Control the **output format** (text, JSON, unformatted markdown, ...)

Refine Results Over Time

It’s a chat!

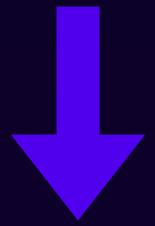

Tell the AI which parts need adjustments

“Remove the Emojis and all hashtags.”

Hands-On!

1. Create a short product announcement text for a new AI-powered website generator
2. Create a Python code snippet that searches & deletes all .png & .jpg files in a given folder
3. Write an email to a colleague that you need feedback on your submitted prototype until end of the week

Key Prompt Engineering Techniques

- Zero- and Few-Shot
- Chain-of-Thought
- Using Delimiters
- Output Templates
- Persona Prompting
- Output Formatting
- Contextual Prompting
- Negative Prompts
- Self-reflective Prompting

Zero- & Few-Shot Prompting

Providing examples can help fine-tune the result tone, style & content

0

Zero-Shot

Provide no examples

“Write a post that explains the core idea behind ChatGPT.”

N

Few-Shot

Provide multiple examples

“Write a post that explains the core idea behind ChatGPT. Use a similar tone & structure as I do in my regular tweets. But don’t use the content. Here are **two example posts**: …”

Finetuning Models

If you need to provide many examples, you could consider finetuning

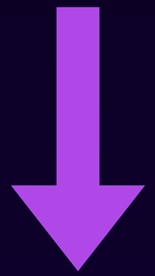

Finetuning **changes the model's internal weights and values**, leading to a different model behavior

Important: Finetuning can be expensive & the model may perform worse at other tasks thereafter!

Include Relevant Information

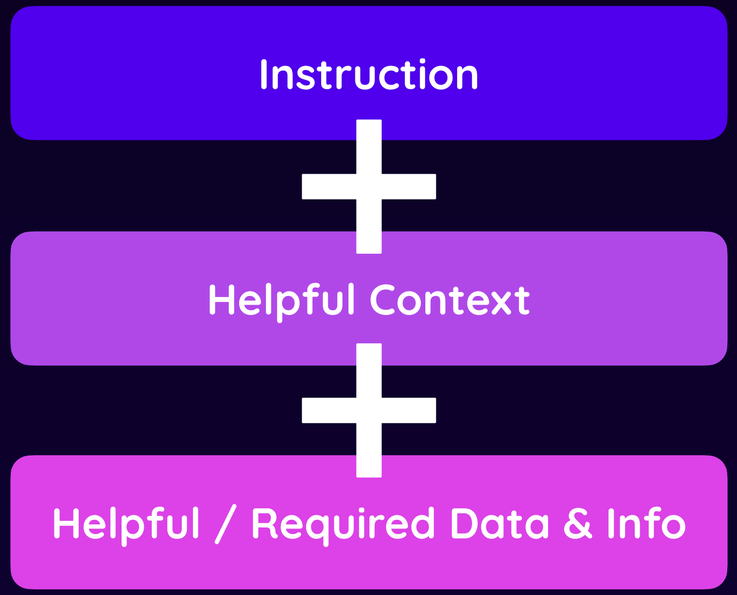

ACADEMIND

Generating Images & Videos with AI

From Text To Images & Video

- Available Options
- Using AI For Generating Images
- Using AI For Generating Videos

AI Image Generation — Available Options

There are many options!

ChatGPT & OpenAI	Other Providers & Models	Integrated Tools
ChatGPT Image Generation	Gemini, Copilot, Grok, ...	Adobe Photoshop
OpenAI API Image Generation	Midjourney	
	Flux, Stable Diffusion	

And many, many others!

ACADEMIND

ChatGPT For Programming

Writing Code with ChatGPT

- Generating Code With & Without Programming Experience
- Debugging & Optimizing Code
- Explaining & Refactoring Code

ChatGPT Is Great For Everyone!

Non-Developers

Build basic programs (utility scripts) & websites with zero development experience

Developers

Boost your productivity, outsource the boring parts, generate dummy data & code faster

ChatGPT For Non-Developers

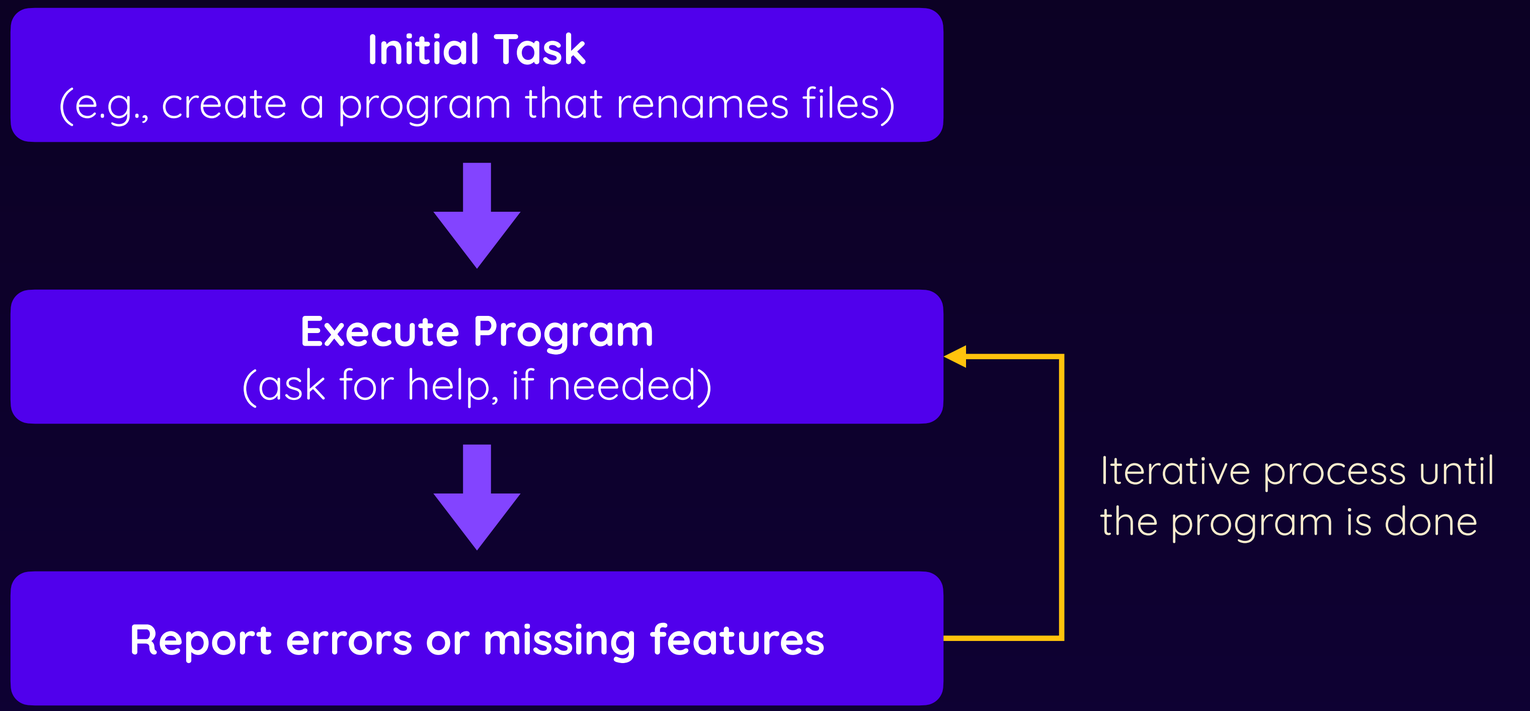

Exercise Time!

Let ChatGPT build a **basic website**!

Starting Page

- Your name
- Image
- List of hobbies

CV Page

- List with career history

Website should have a modern, clean, dark-mode styling.

ACADEMIND

Be Careful

Code you don't know could cause harm!

Depending on your request & prompt it could delete files, erase data, crash your system etc.

ChatGPT 🤝 Developers

As a developer, you can massively **boost your productivity** by using ChatGPT!

ChatGPT For Developers

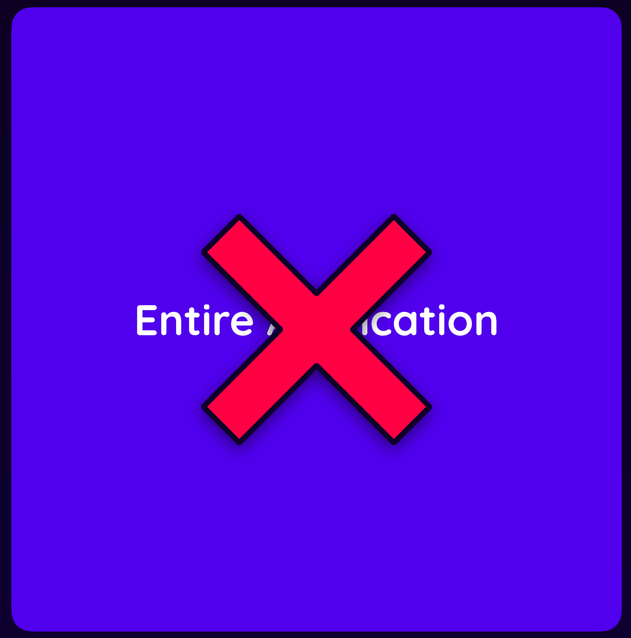

Don’t ask ChatGPT for entire applications or

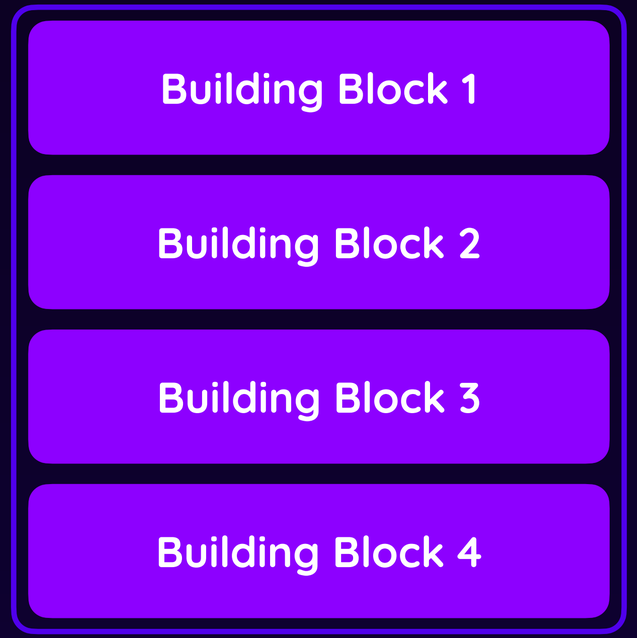

Combine manually (or with additional help from ChatGPT)

Instead, prefer asking for individual building blocks

ChatGPT For Developers

Building Blocks > Entire Apps

Use ChatGPT to speed up the development of the individual application building blocks

Iterative Development

Add more and more features by splitting your requests across multiple prompts

Refine Code Manually

Instead of deriving fancy prompts, consider performing fine-tuning tasks manually

Use ChatGPT For Debugging

Report errors & bugs (+ relevant code snippets) to ChatGPT to speed up debugging

Explain Code

Instruct ChatGPT to quickly explain & summarize unfamiliar code

Use ChatGPT for Refactoring

Let ChatGPT refactor code or use ChatGPT to get improvement ideas

Don’t Limit Yourself To Just ChatGPT!

ChatGPT

Great for generating entire building blocks

+

GitHub Copilot

Great for generating smaller code snippets “on the fly” & fine-tuning your code

Useful Prompting Techniques

If Code Gets Cut Off

“Output ‘Continuing’ and continue”

Skip Explanations

“Provide just the code without any extra explanations or text.”

Add Context To Errors

“The user authentication code seems to break the program with the following error message: [message]”

Let ChatGPT Improve Itself

“How could the code be improved?”

Generate Dummy Data with ChatGPT

You’re not limited to generating code

Generate dummy data with ChatGPT
(e.g., dummy users)

Hands-On: ChatGPT Content Creation

Practicing How To Generate Content With ChatGPT

- Create & Advertise a Realistic Blog Post
- SEO
- Add Images (Midjourney)
- Creating a Video Script

Creating Content with ChatGPT

Plan a Blog Post
Keywords + Outline

Write the Blog Post

Add Images
(incl. Thumbnail)

SEO

Format as Markdown

Create YouTube Video Script
(based on Blog Post)

Create Newsletter Email

Create Marketing Tweets

Topic

Understanding ChatGPT & How It Works Behind The Scenes

ACADEMIND

Hands-On: Programming & ChatGPT

Practicing How To Write Code With ChatGPT

- Non-Developer Example: Building a “Monster Slayer” Game
- Developer Example: Building a Meetups REST API

Building a “Monster Slayer” Game

Using ChatGPT with **no / minimal** programming experience

Building a **text-based / command-line based** game: “Monster Slayer”

Description

It’s a **turn-based** game where the user (= player) fights a monster (= computer).
During every turn, the player can perform a **regular** or **strong attack** or **heal**.
The **strong attack** should only be available **every three turns**. **Healing** should only be available **every five turns**.
After each turn, the **monster attacks**.
Damage & heal values are calculated **randomly**.
The first participant to go below **0 health loses**.
Both participants **start with 100 health**.
Once the game is over, the **winner** should be **displayed on the screen** and the player should be **asked if a new game** should be started.

Enhancing The “Monster Slayer” Game

Add Username

Allow the user to choose a username when the program starts

Add Difficulty Levels

Adjust damage & heal values based on chosen level

Manage High Score

Save the number of required turns in text file

Display the current high score after every game

ACADEMIND

Results May Deviate

The Code You Get Will Likely Differ

Because of built-in randomness, different prompts & GPT model evolution, you will very likely not get the same results!

Building a “Meetups” REST API

Using ChatGPT **with** programming experience

Building a **NodeJS REST API**

Description

It’s a REST API with the following **endpoints**:
POST /meetups → Create a new meetup
GET /meetups → Fetch meetups
PATCH /meetups/<id> → Update existing meetup
DELETE /meetups/<id> → Delete existing meetup

Every meetup has an **id, title, summary & address**.
Meetup data should be **stored** in a **meetups.json** file, incoming data must be **validated**.
Data should be exchanged in **JSON format.**

ACADEMIND

Using AI Programmatically via APIs

Building AI-powered Applications

- ChatGPT vs OpenAI API
- Using the OpenAI API

API
Application Programming Interface

A set of rules and protocols that allow software
to communicate with other software

For example, the OpenAI API allows your software
(e.g., some mobile app, or some internal tool) to use
OpenAI's AI models through code

ChatGPT vs GPT APIs

AI Chatbots

e.g., ChatGPT, Gemini, ...

AI-powered applications built by AI companies (which may be the same companies that built the underlying AI models)

Used by "end users" to interact & solve problems

AI APIs

e.g., OpenAI API, Gemini Dev API, ...

Programmatic access to AI models (often — but not always — provided by the companies that developed the models)

Used by developers to build their own AI-powered apps

Prerequisites

- Interested in building AI-powered apps
- Programming Experience
- Python (or some other language)

ACADEMIND

RAG, CAG & Finetuning

Expanding The Knowledge of AI Models

- Understanding RAG & CAG
- Building RAG & CAG Workflows
- Finetuning

Problem

Key data may be unknown to the AI model

E.g., personal data, company-internal data, or data generated after the knowledge cutoff date

Making Sense of RAG & CAG

RAG

Retrieval-Augmented Generation

CAG

Cache-Augmented Generation

Both are about enhancing the AI prompt with extra information that allows the AI model to generate a better response

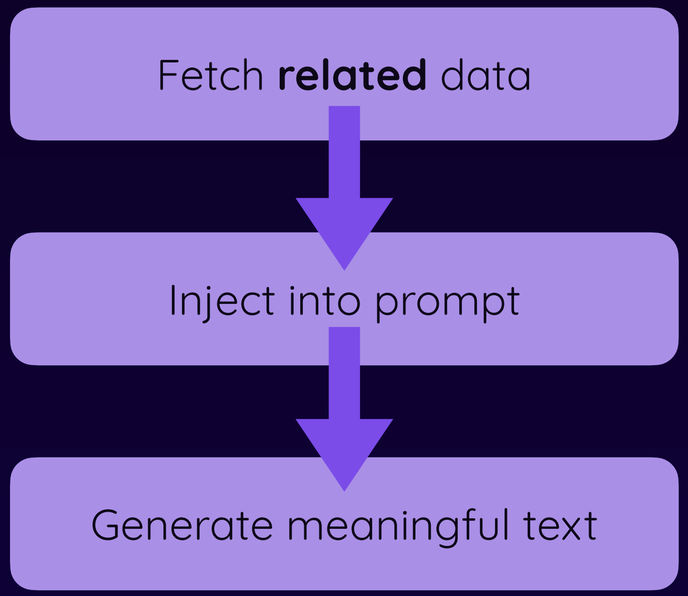

Efficient, secure but requires more complex setup & may miss loosely related data

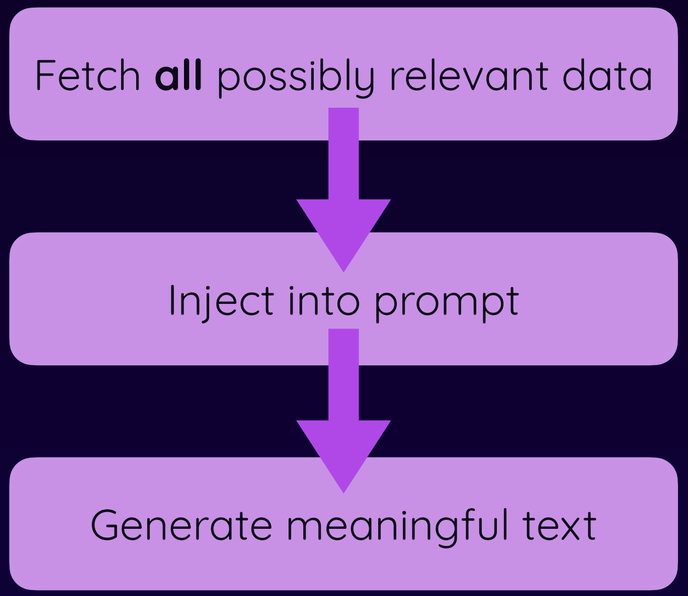

Less efficient, requires large context window but is less complex & can help with loosely related data

Building CAG Workflows

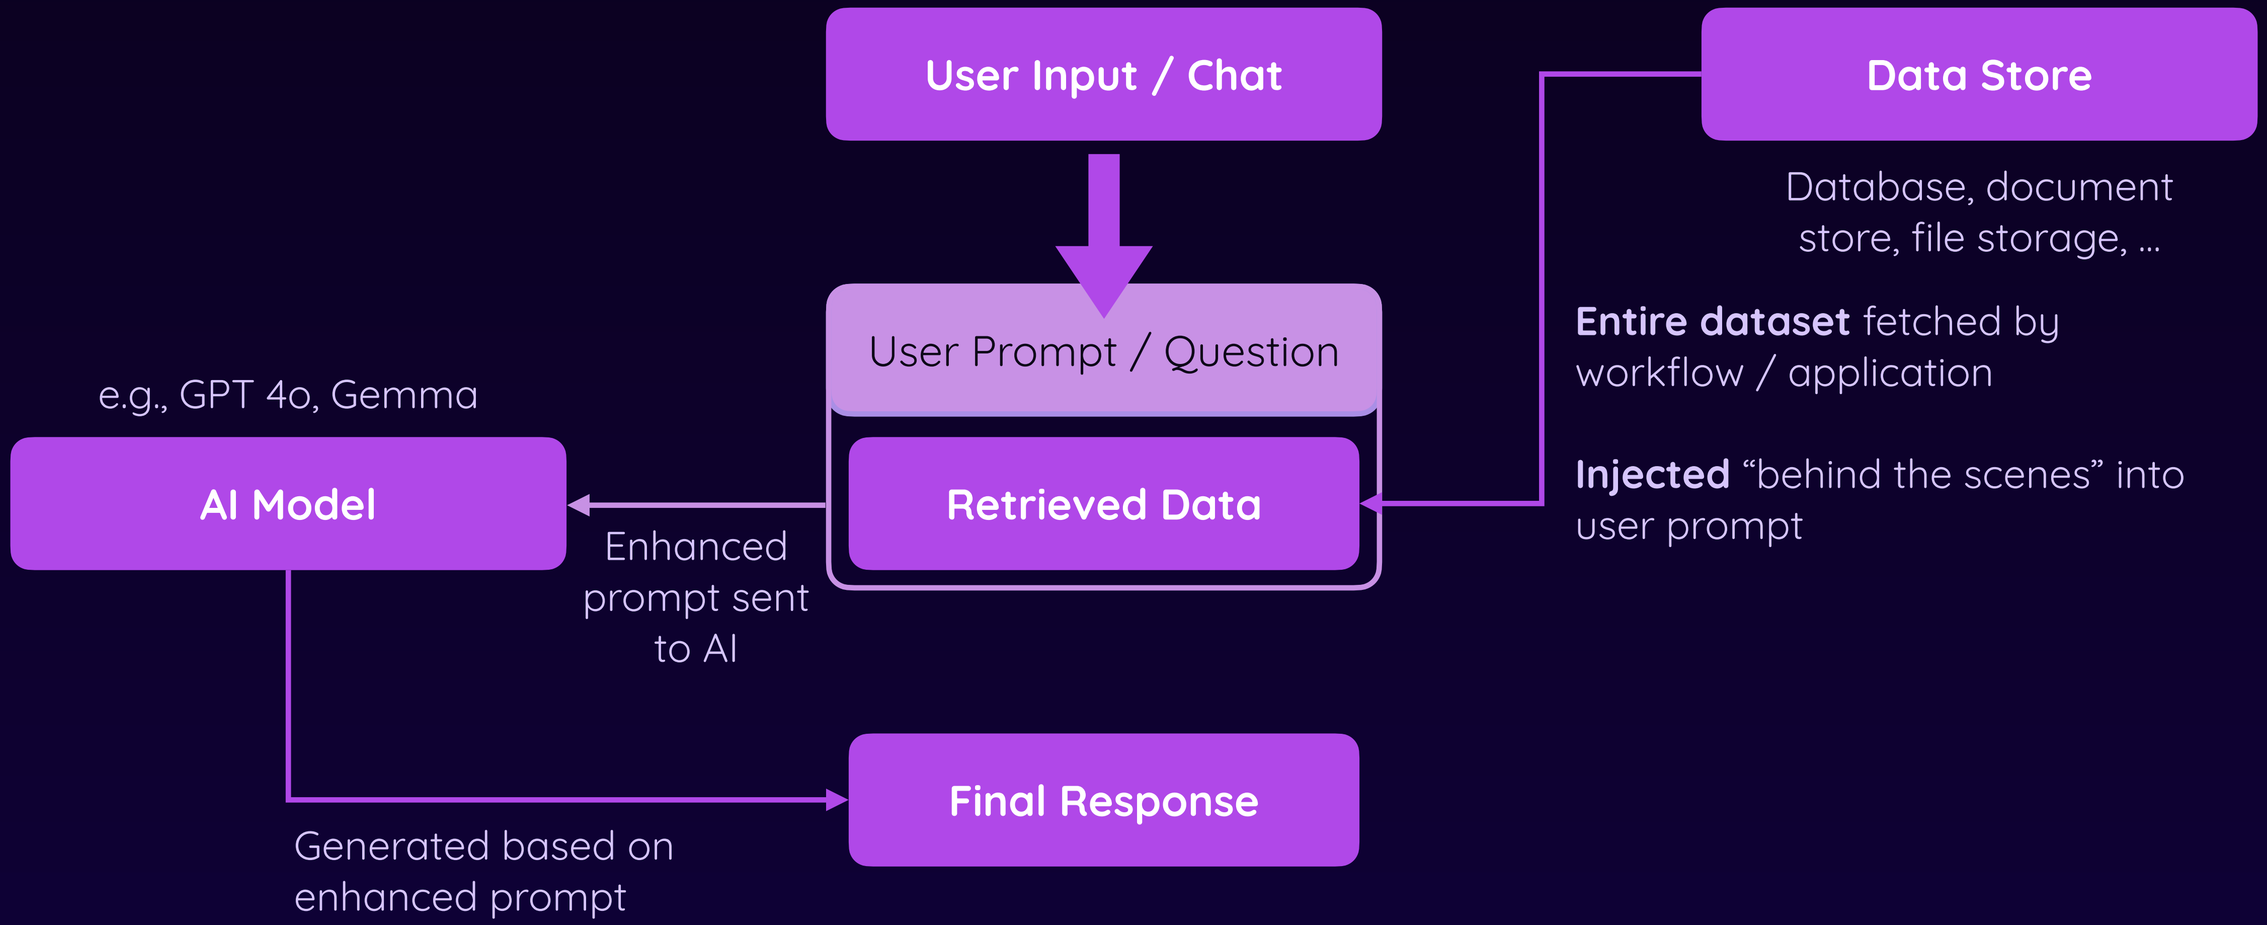

Building RAG Workflows

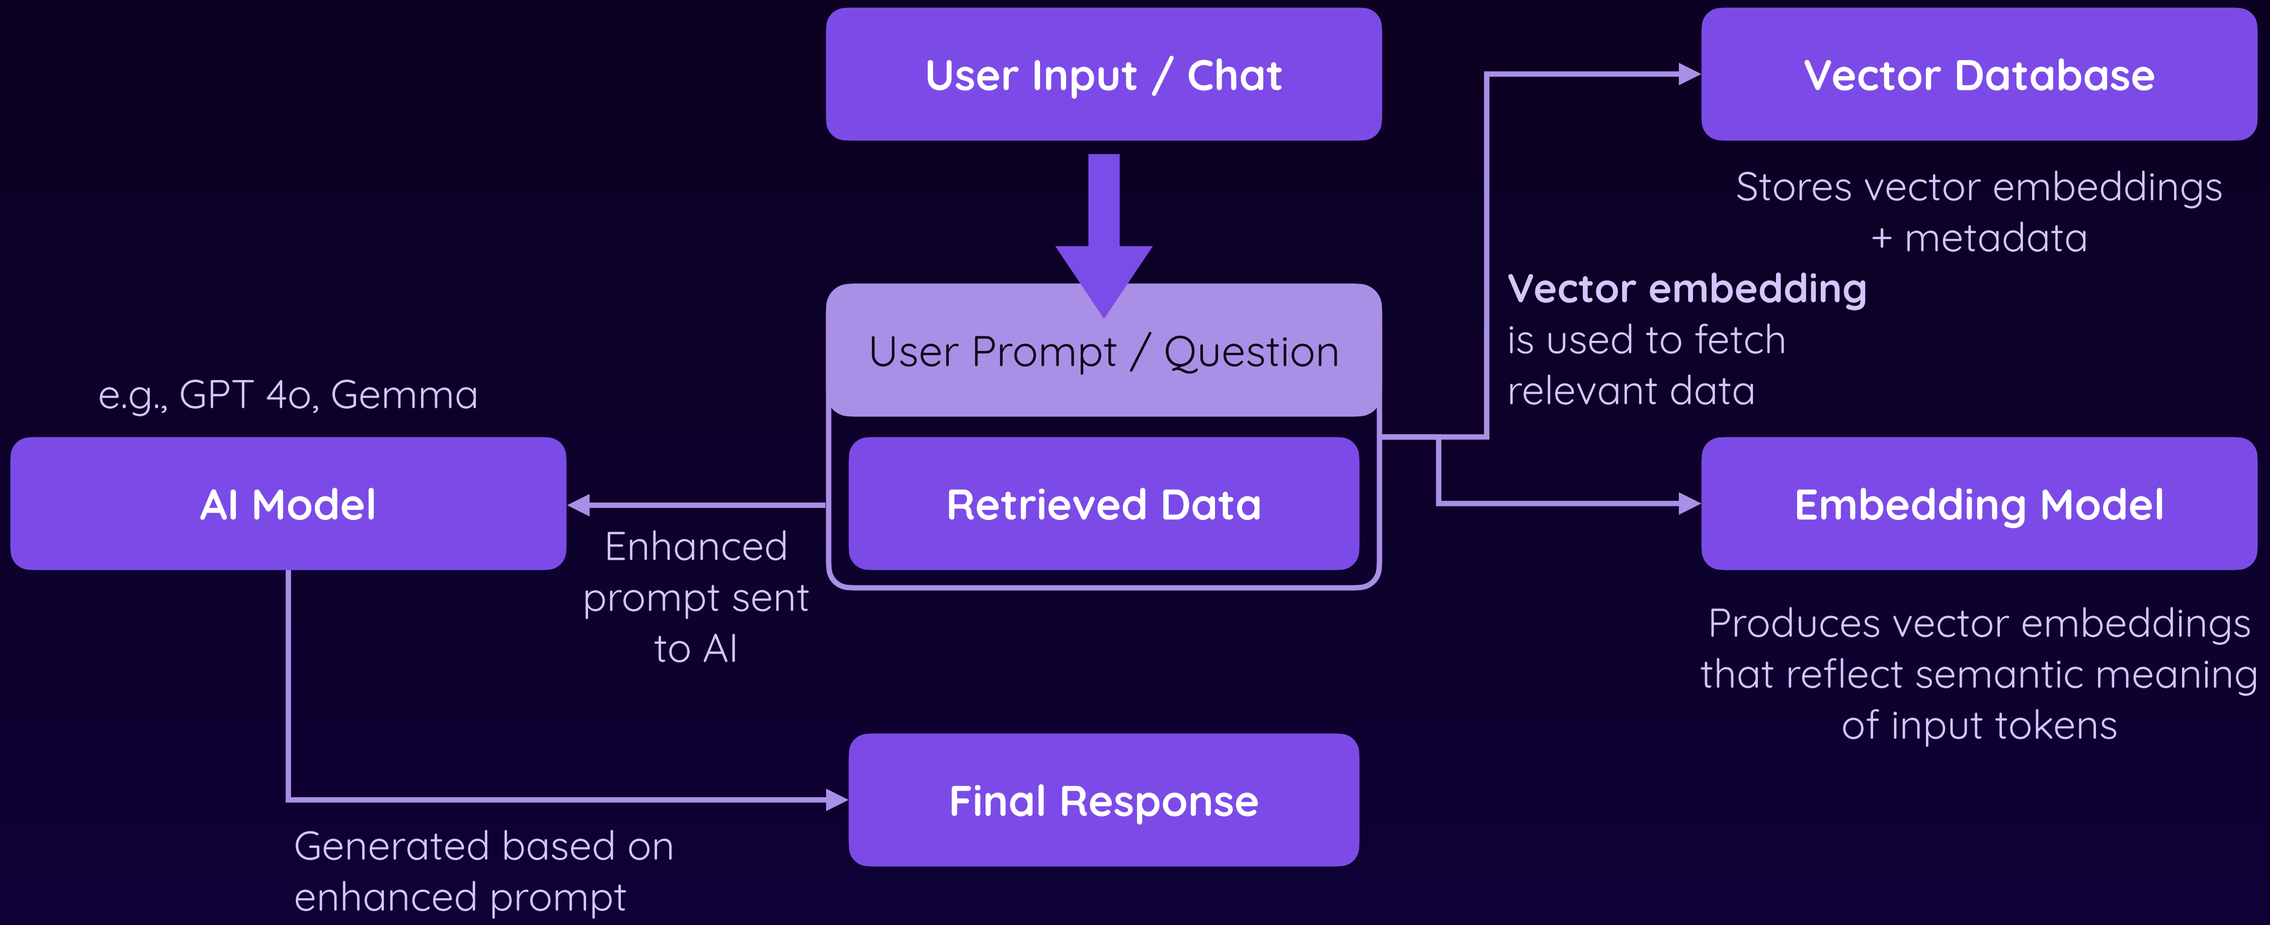

Vector Embeddings Represent Relations

Related tokens / words are stored in similar places

Important: In reality, it's a n-dimensional space!

RAG / CAG & Few-Shot Prompting

RAG / CAG

Data / examples are inserted into prompt

Without the enhanced prompt, the model won't know about the data

Few-Shot Prompting

Examples are inserted into prompt

Without the enhanced prompt, the model won't know about the examples

RAG / CAG is a form of few-shot prompting

Typically an automated process to enrich prompts with data

Typically a manual process to enrich prompts with examples

Finetuning Models

Unlike few-shot prompting or RAG / CAG, Finetuning is about **permanently adjusting** the weights (and therefore **"knowledge" and behavior**) of the underlying model

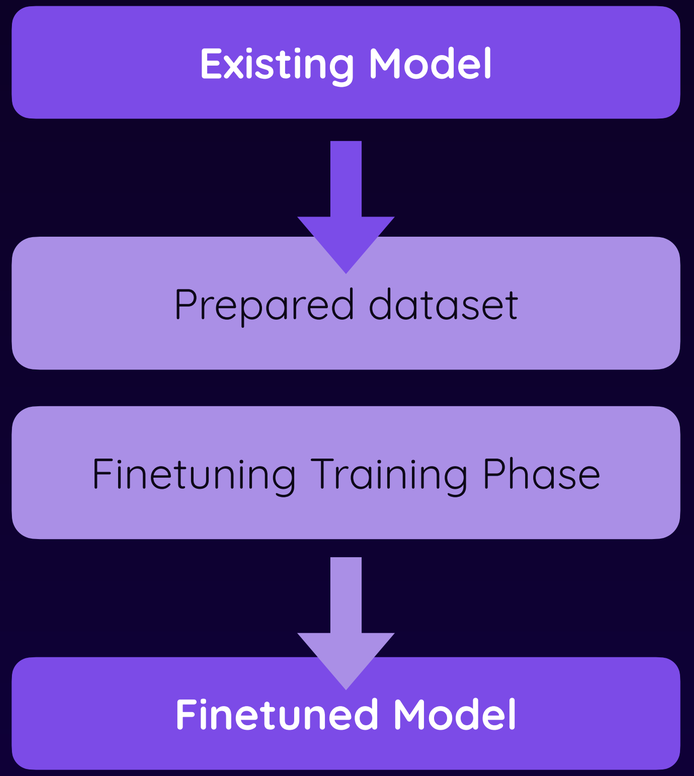

Doesn't know anything about your data

E.g., selected list of posts or company-internal documents

Can take hours or even days
Potentially very high cost!

Knows the finetuning data and related behaviors (e.g., writing style)

Finetuning vs RAG / CAG

RAG / CAG / Few-shot	Finetuning
Data is inserted into prompt	Existing model is re-trained on custom data
Without the enhanced prompt, the model won't know about the data	Model "learns" about data & specific behaviors / text style
Low / mediocre complexity	Mediocre / high complexity
Cost may increase for long prompts / lots of data	Mediocre / high training cost BUT potentially lower prompting costs
Prefer in most use-cases, especially when context window size & data fetching is not an issue	Prefer if you need a specialized model or when data fetching is not an option

ACADEMIND

Automation & Agents

Building Automated AI-powered Workflows

- AI Automation vs AI Agents
- Building Agents Without Writing Code
- Building Agents With Code

Automation vs Agents

AI Automation

A pre-defined workflow that uses AI in one or more steps

AI Agents

An AI model that can solve a variety of tasks by using assigned tools

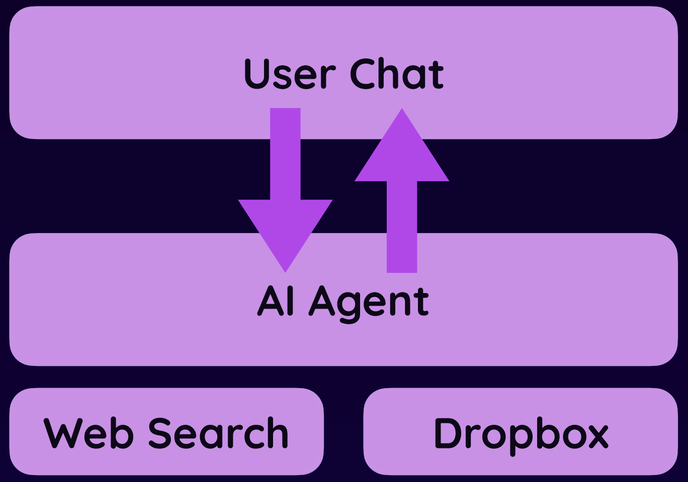

Building AI Workflows (Automation & Agents)

You got many options & platforms

No Code	With Code
n8n	Custom code
Gumloop	LangChain & LangGraph
Flowise	Pydantic AI
Trilex AI	CrewAI
And many, many others...	Agno
	Vercel AI SDK
	And many, many others...

Enhancing The “Meetups” REST API

Add Authentication

Add POST /signup & POST / login routes

Implement JWT-based authentication

Protect all routes except for GET /meetups

Handle Errors

Throw errors & use generic error handling middleware

Use appropriate error status codes (e.g., 401 if not authenticated)

ACADEMIND

Developer AI Tools

Beyond ChatGPT

- GitHub Copilot & GitHub Copilot Chat
- Cursor IDE

ACADEMIND

You’re Sharing Data!

Use these AI tools with care

Carefully evaluate the vendors’ privacy statements & consider if you’re allowed (and want to) share code & prompts!

AI Tools Covered

GitHub Copilot	GitHub Copilot Chat	Cursor IDE
Smart code completions	AI chat interface integrated in IDE	IDE based on VS Code with integrated AI
Enhance the default IDE auto-completions	Uses project code as context	Prompt-focused, AI-driven programming
Can be triggered in different ways & requires no prompt writing	Use for explanations, fixes, test generation & more	Generate, explain & enhance code, fix errors, search docs & more
Paid Plans **General Availability**	**Copilot Extension** **Limited Availability**	**Free & Paid Plans** **General Availability**

ACADEMIND

Building Automated AI Workflows

Just Define The Goal, Not The Steps

- Get Started with AutoGPT
- Explore LangChain For Custom AI Workflows

What Is AutoGPT?

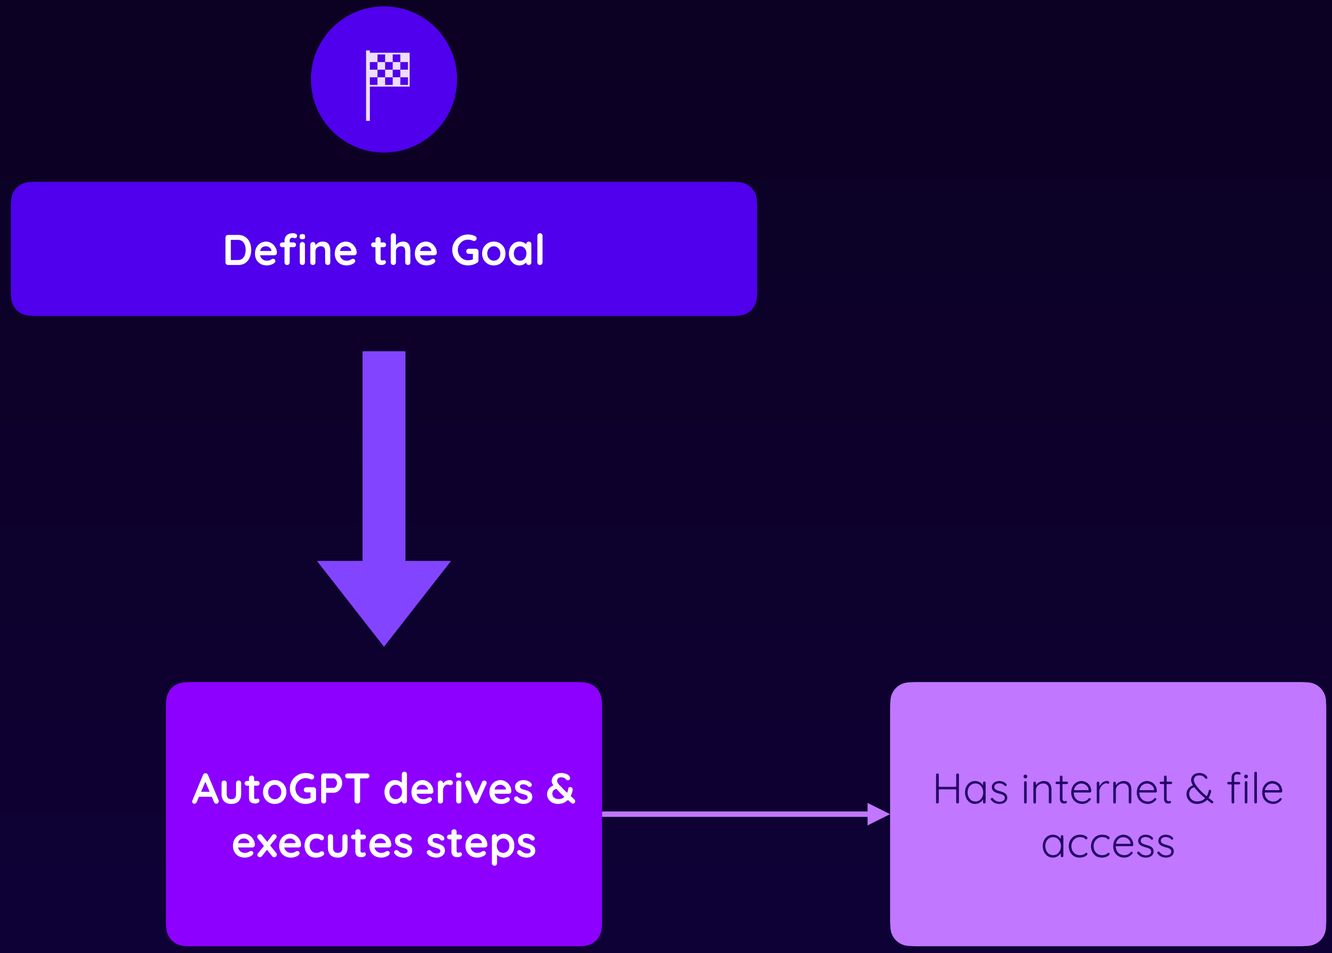

AutoGPT Can Be Awesome

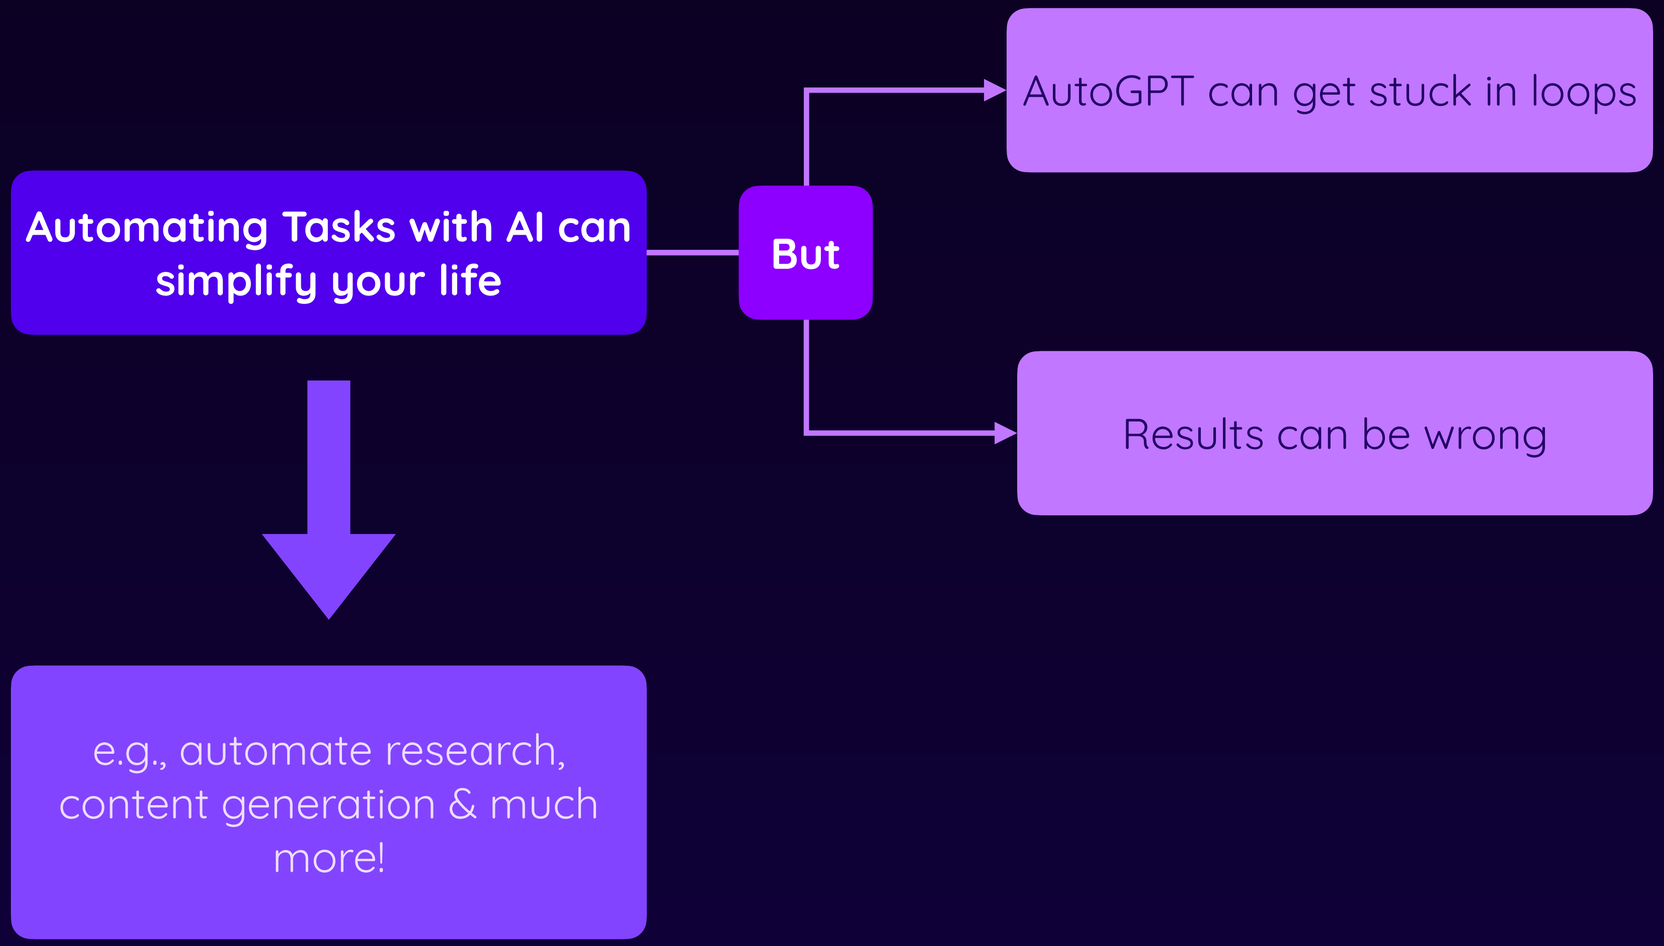

What Is LangChain?

A Framework For Building Your Own Automated AI Workflows

Requires coding skills!

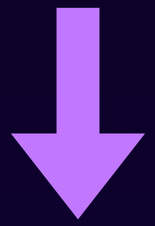

Create any kind of AI workflow

Build an email generator

Build a web research tool

Build anything!